

COM 3240

# REINFORCEMENT LEARNING



# OPTIMISATION

## GRADIENT DESCENT

**Desirable:** A system (an agent) that performs well a task

The system has parameters that we need to select appropriately (optimise)

$f(x_1, x_2, \dots, x_n)$   
parameters

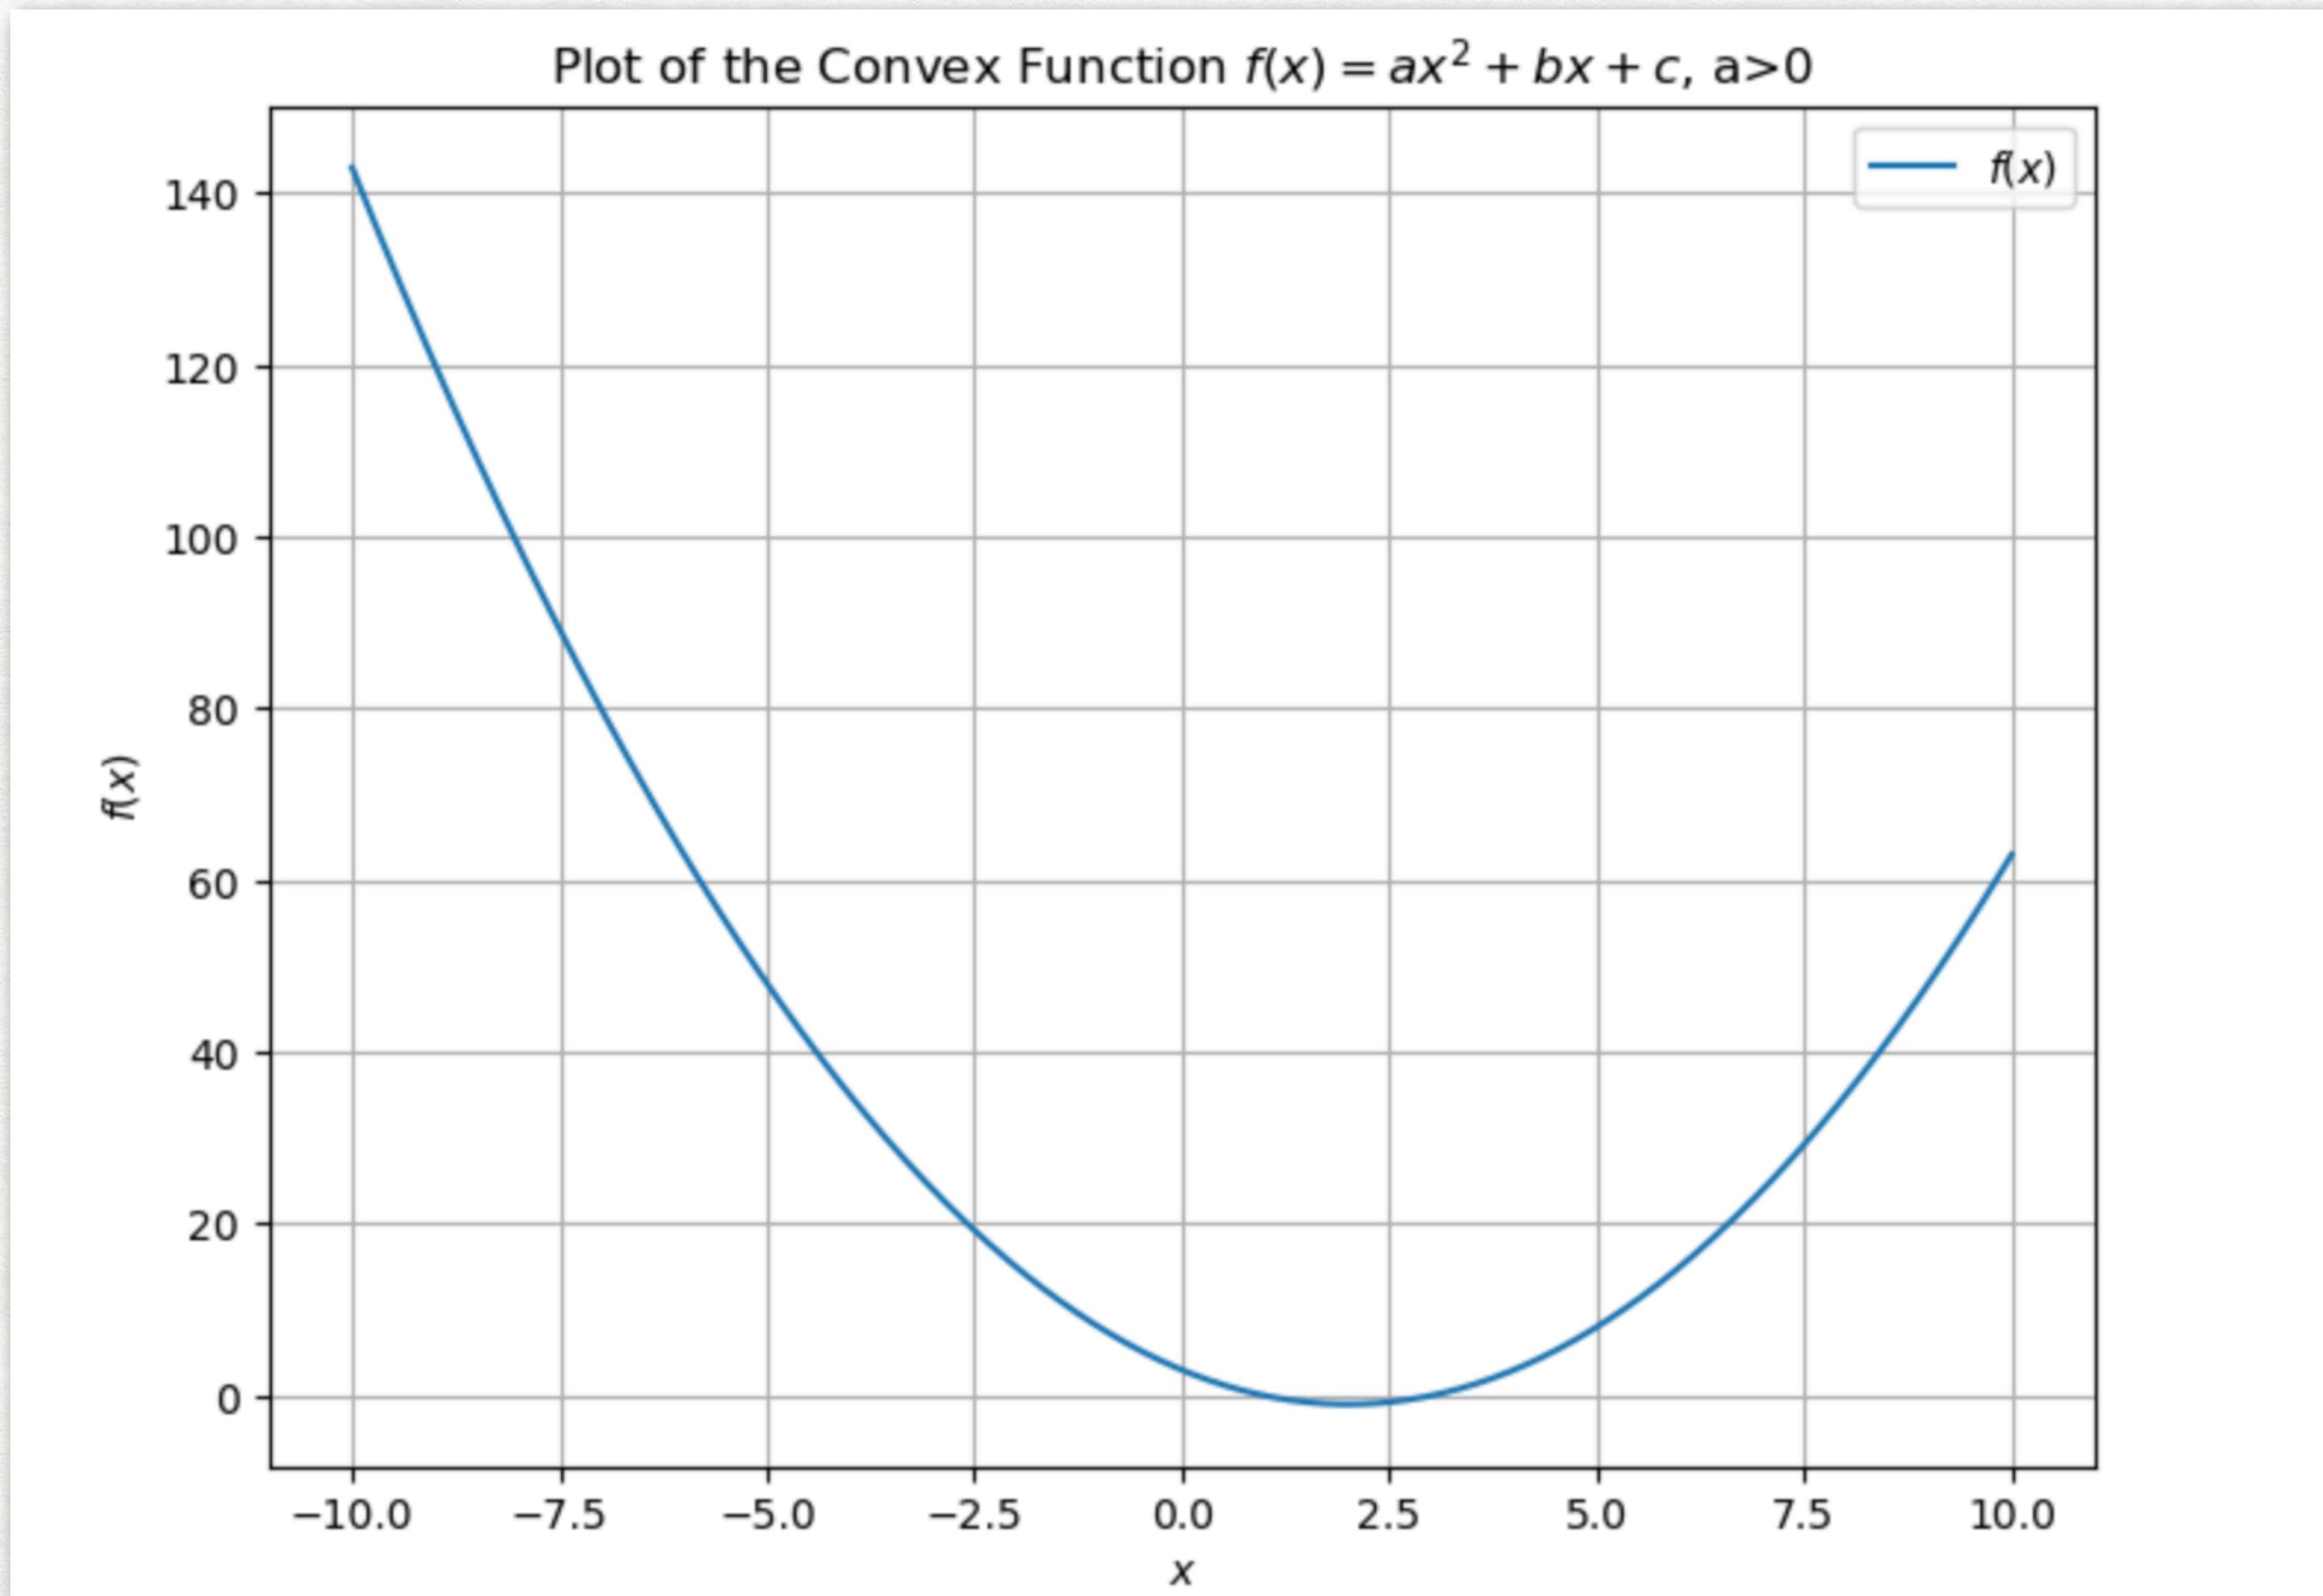
$f$  System's performance: maximise  
System's error : minimise



DALL·E

# OPTIMISATION

## SINGLE (GLOBAL) MINIMUM

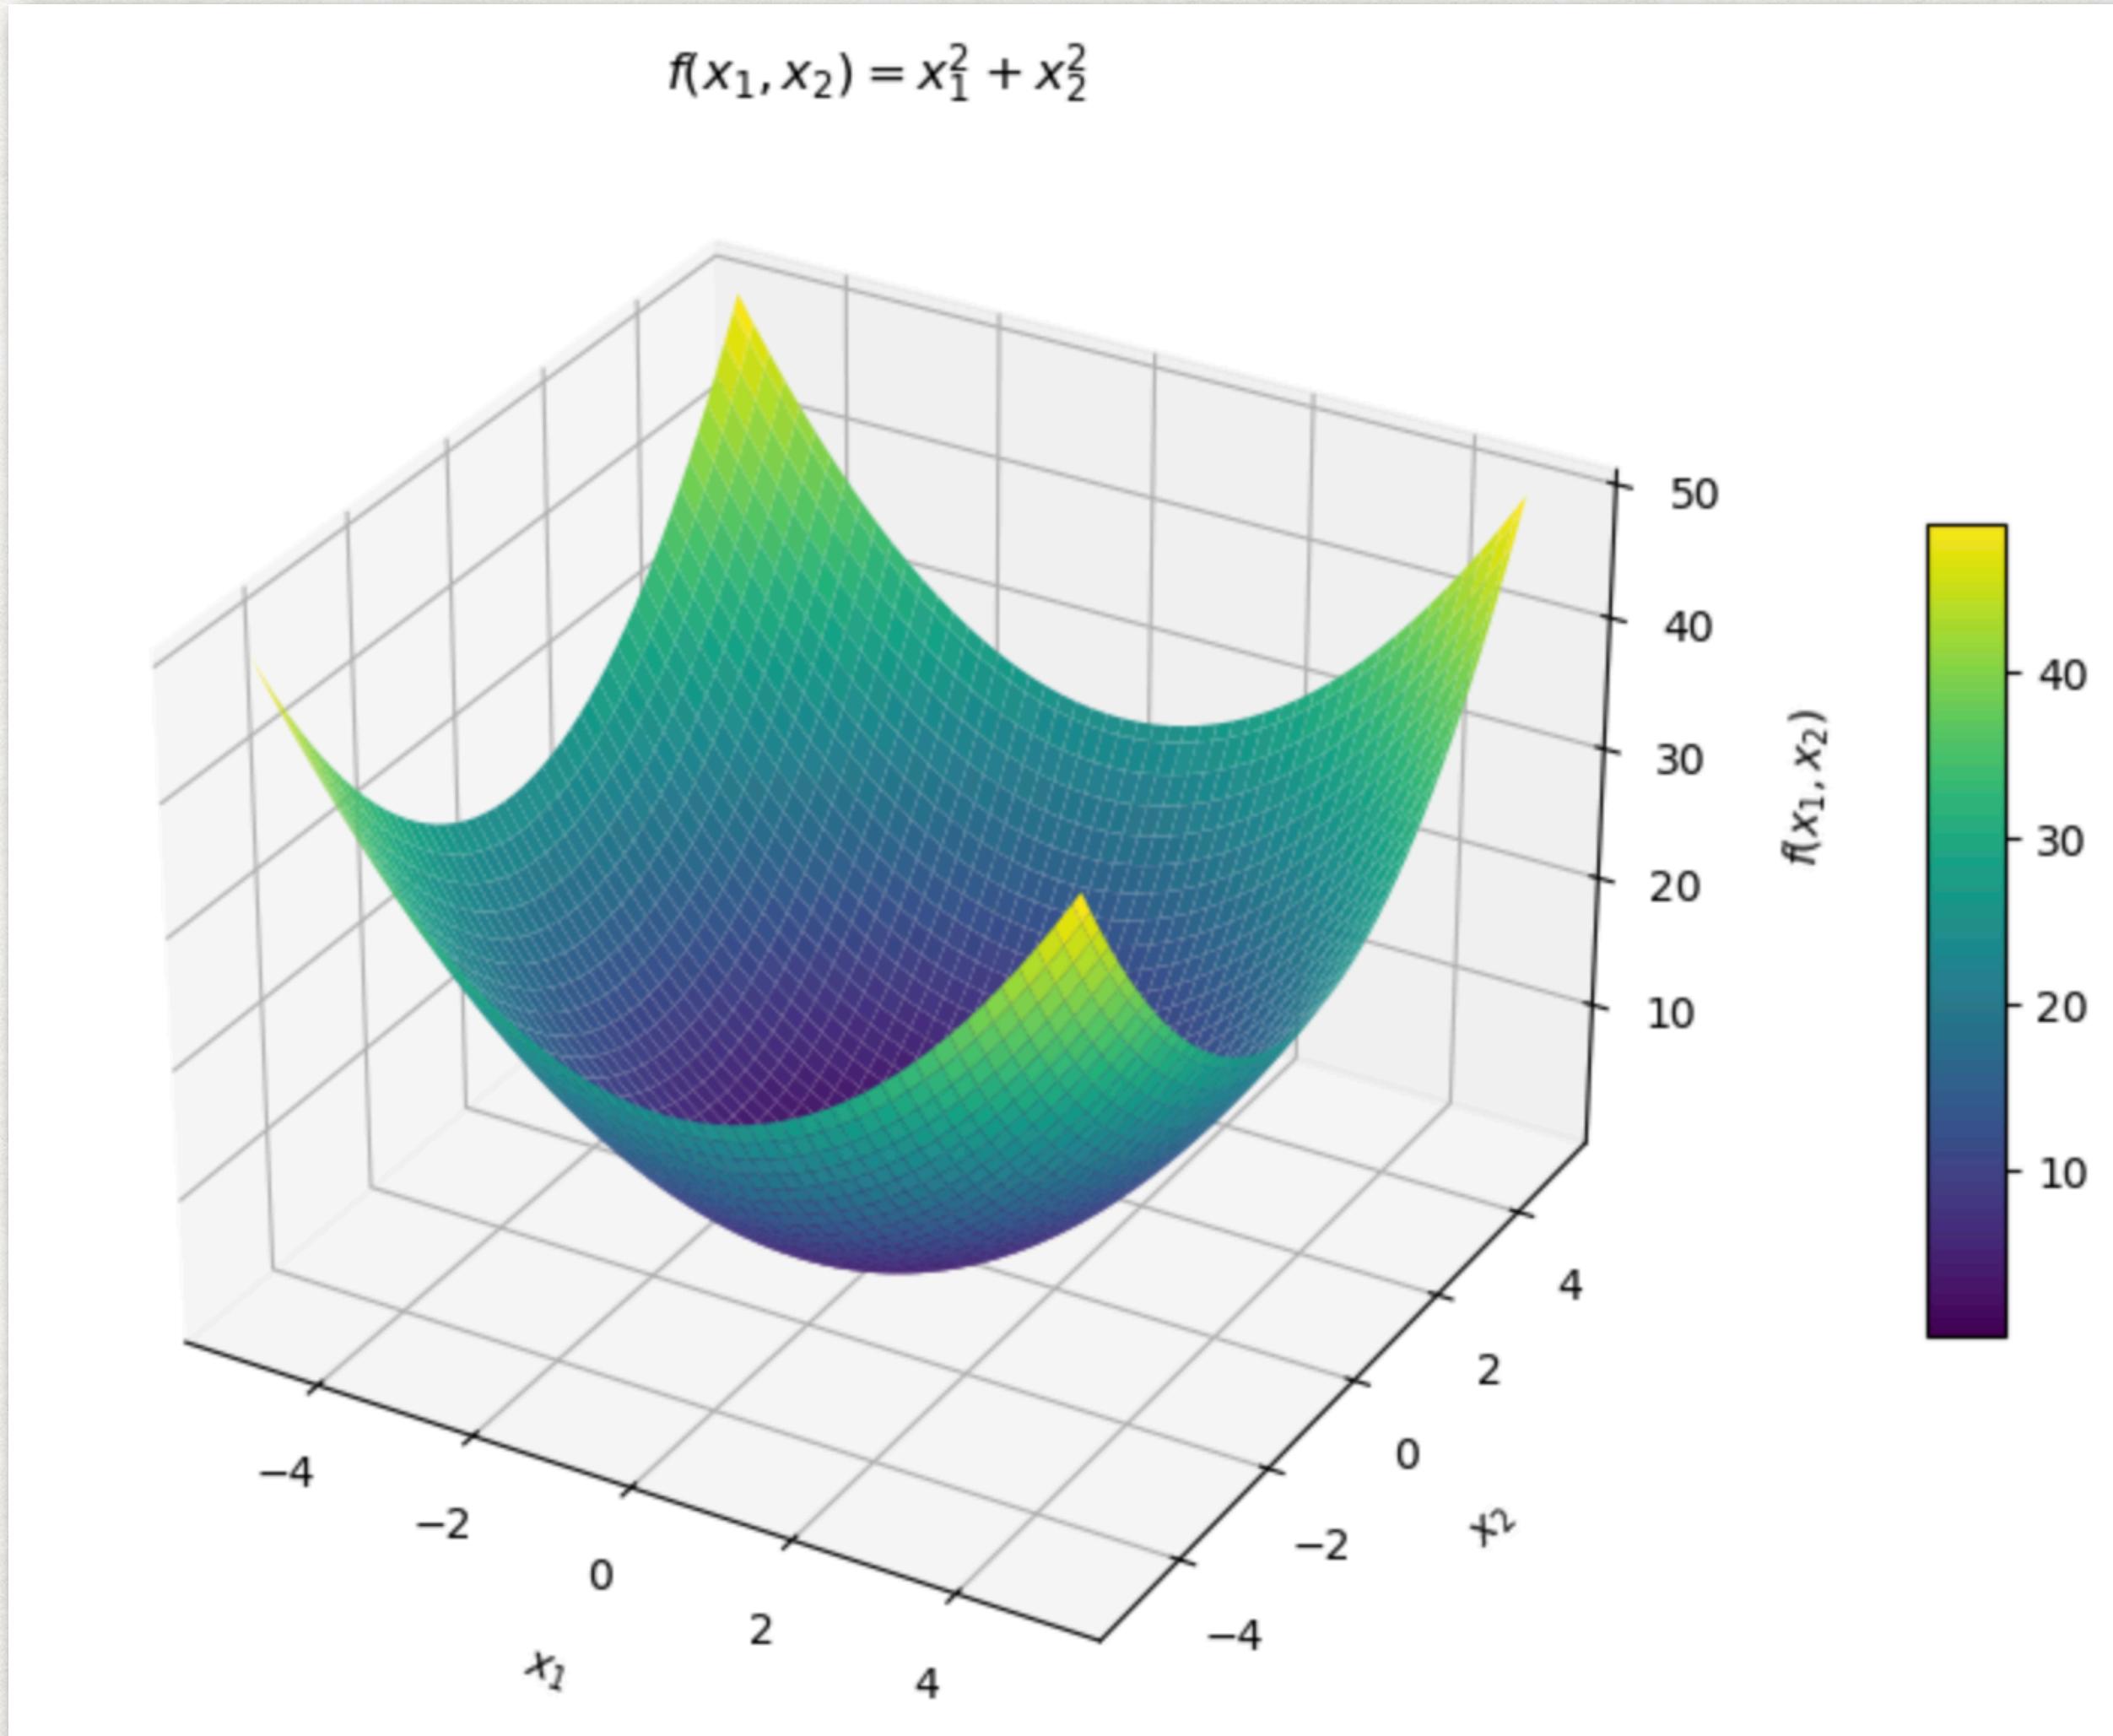


$$\frac{d}{dx}(ax^2 + bx + c) = 2ax + b$$

$$2ax + b = 0$$

# OPTIMISATION

## SINGLE (GLOBAL) MINIMUM

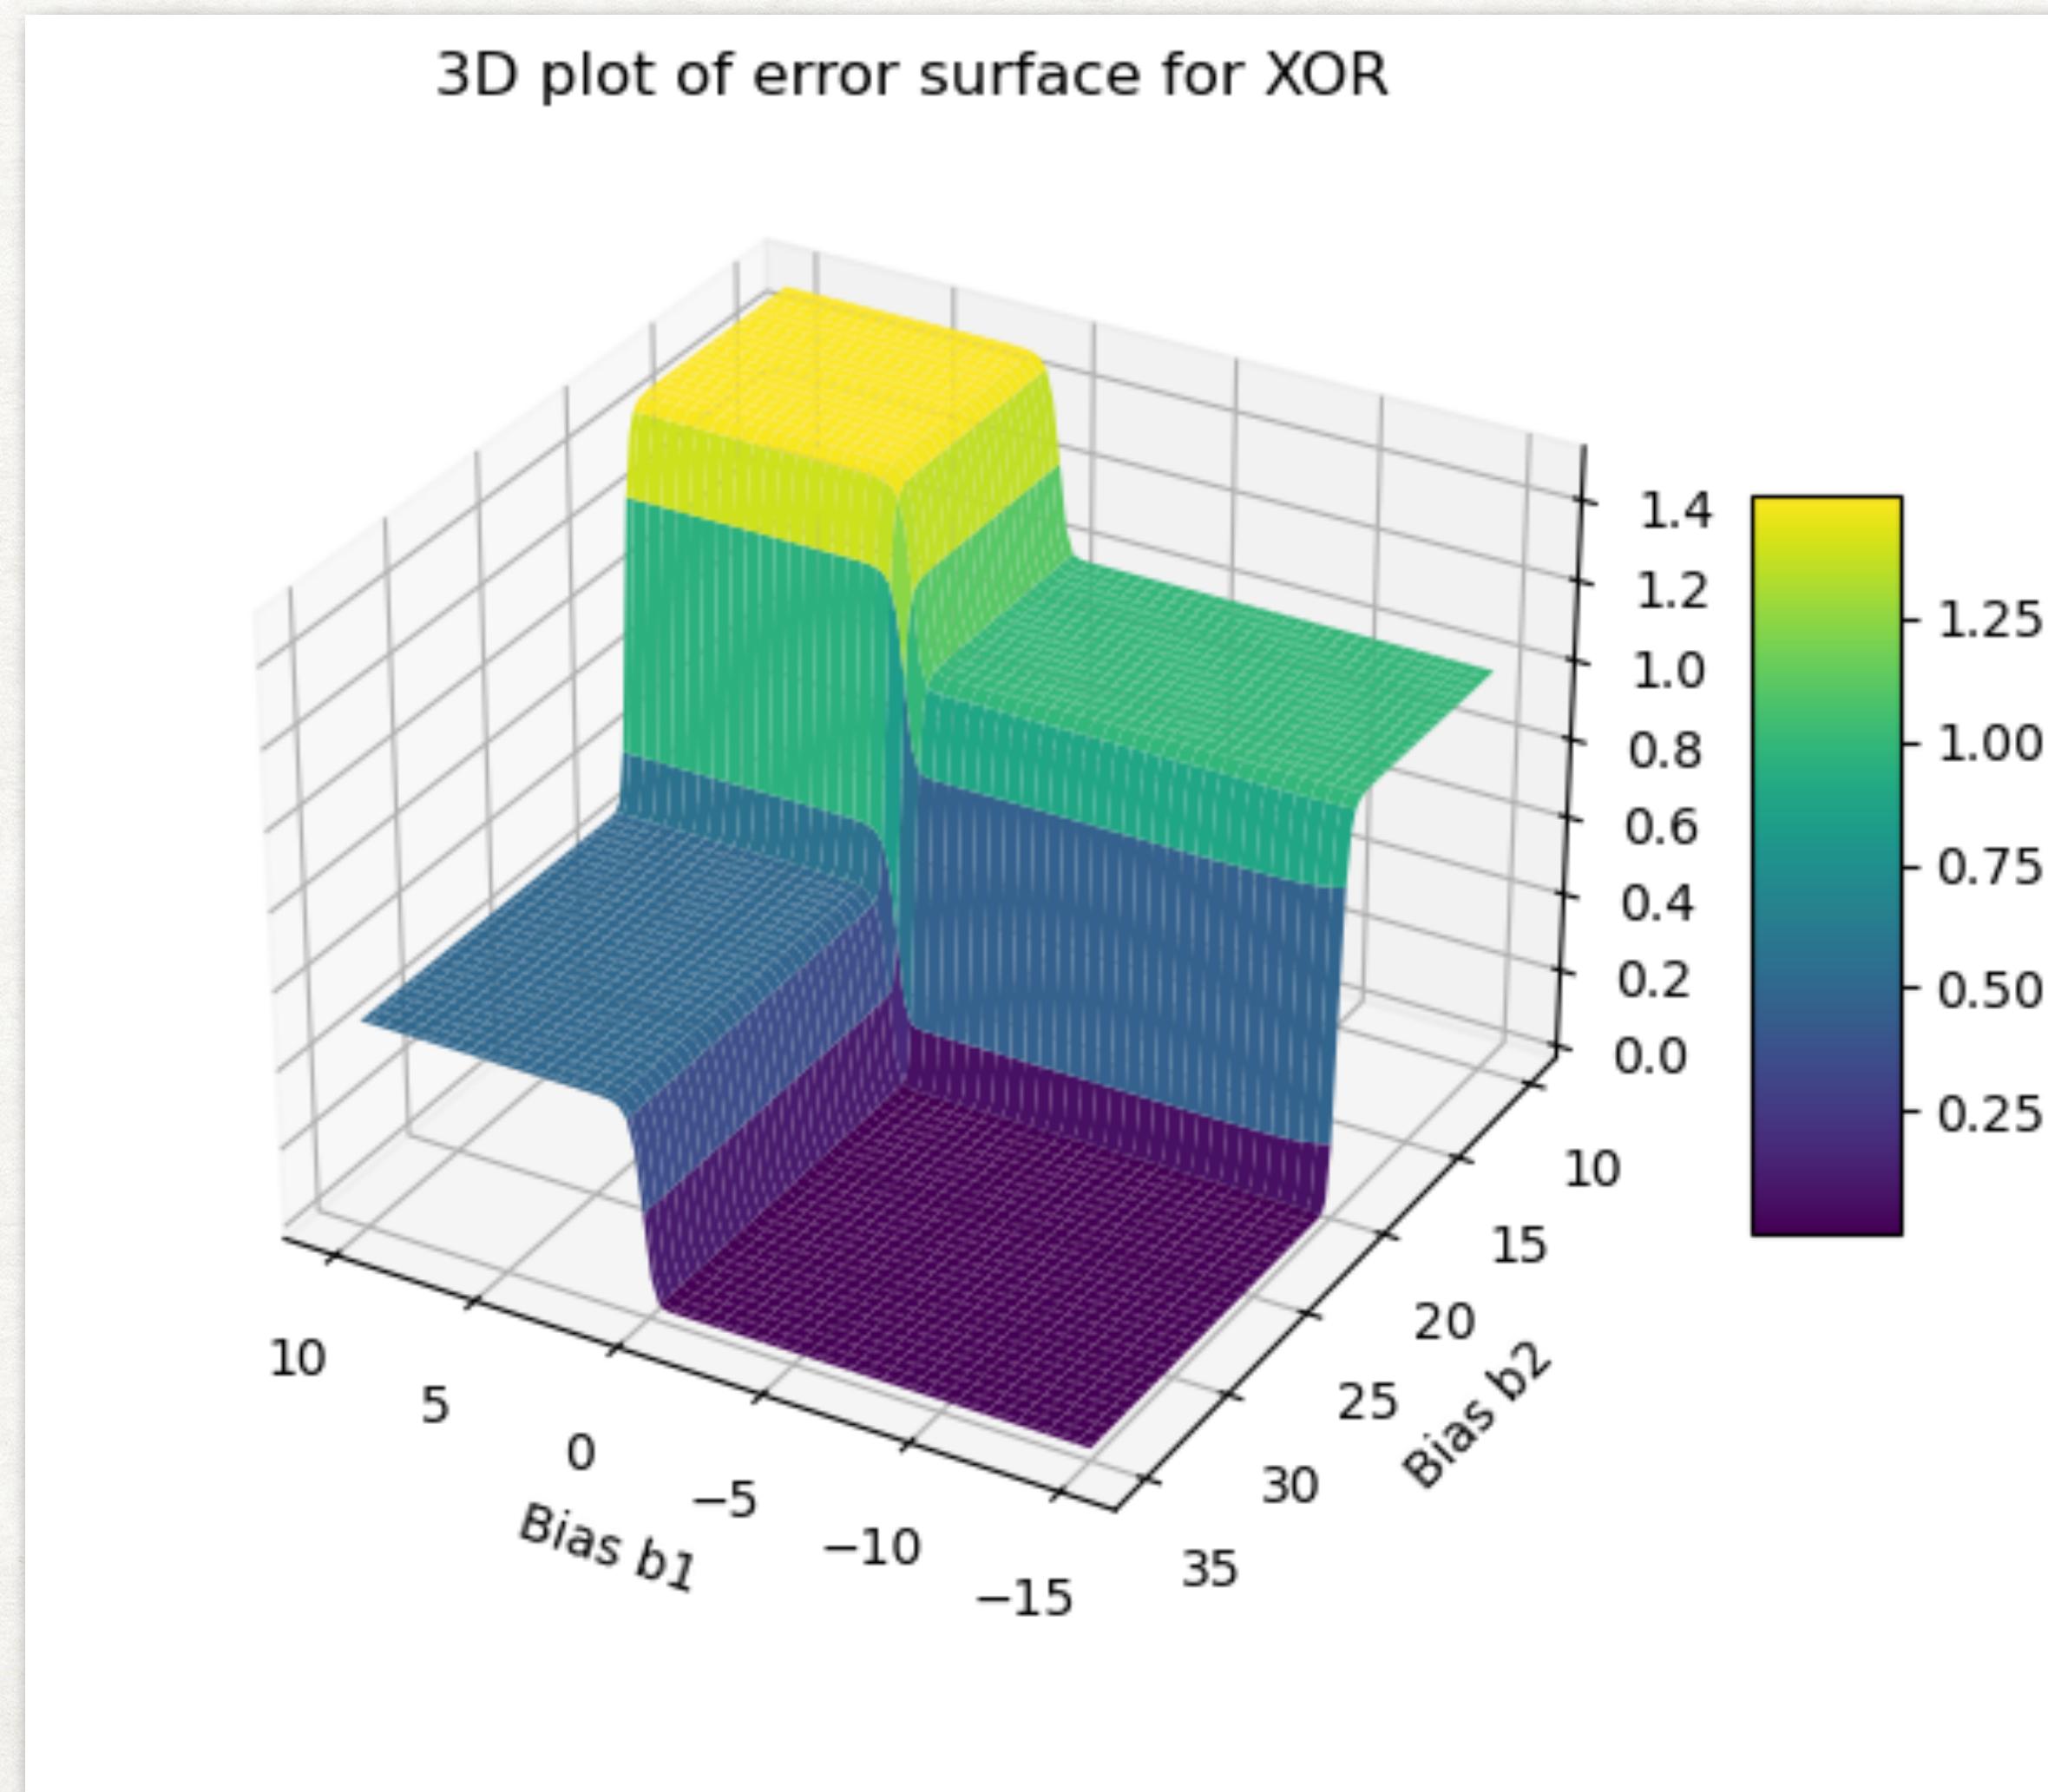


$$\nabla f(\mathbf{x}) = \left[ \frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2} \right] = 0$$

**But this is not the case for all machine learning techniques, including multi-layer and deep neural networks**

# OPTIMISATION - GRADIENT DESCENT

## MULTIPLE LOCAL MINIMA



Simple ANN example

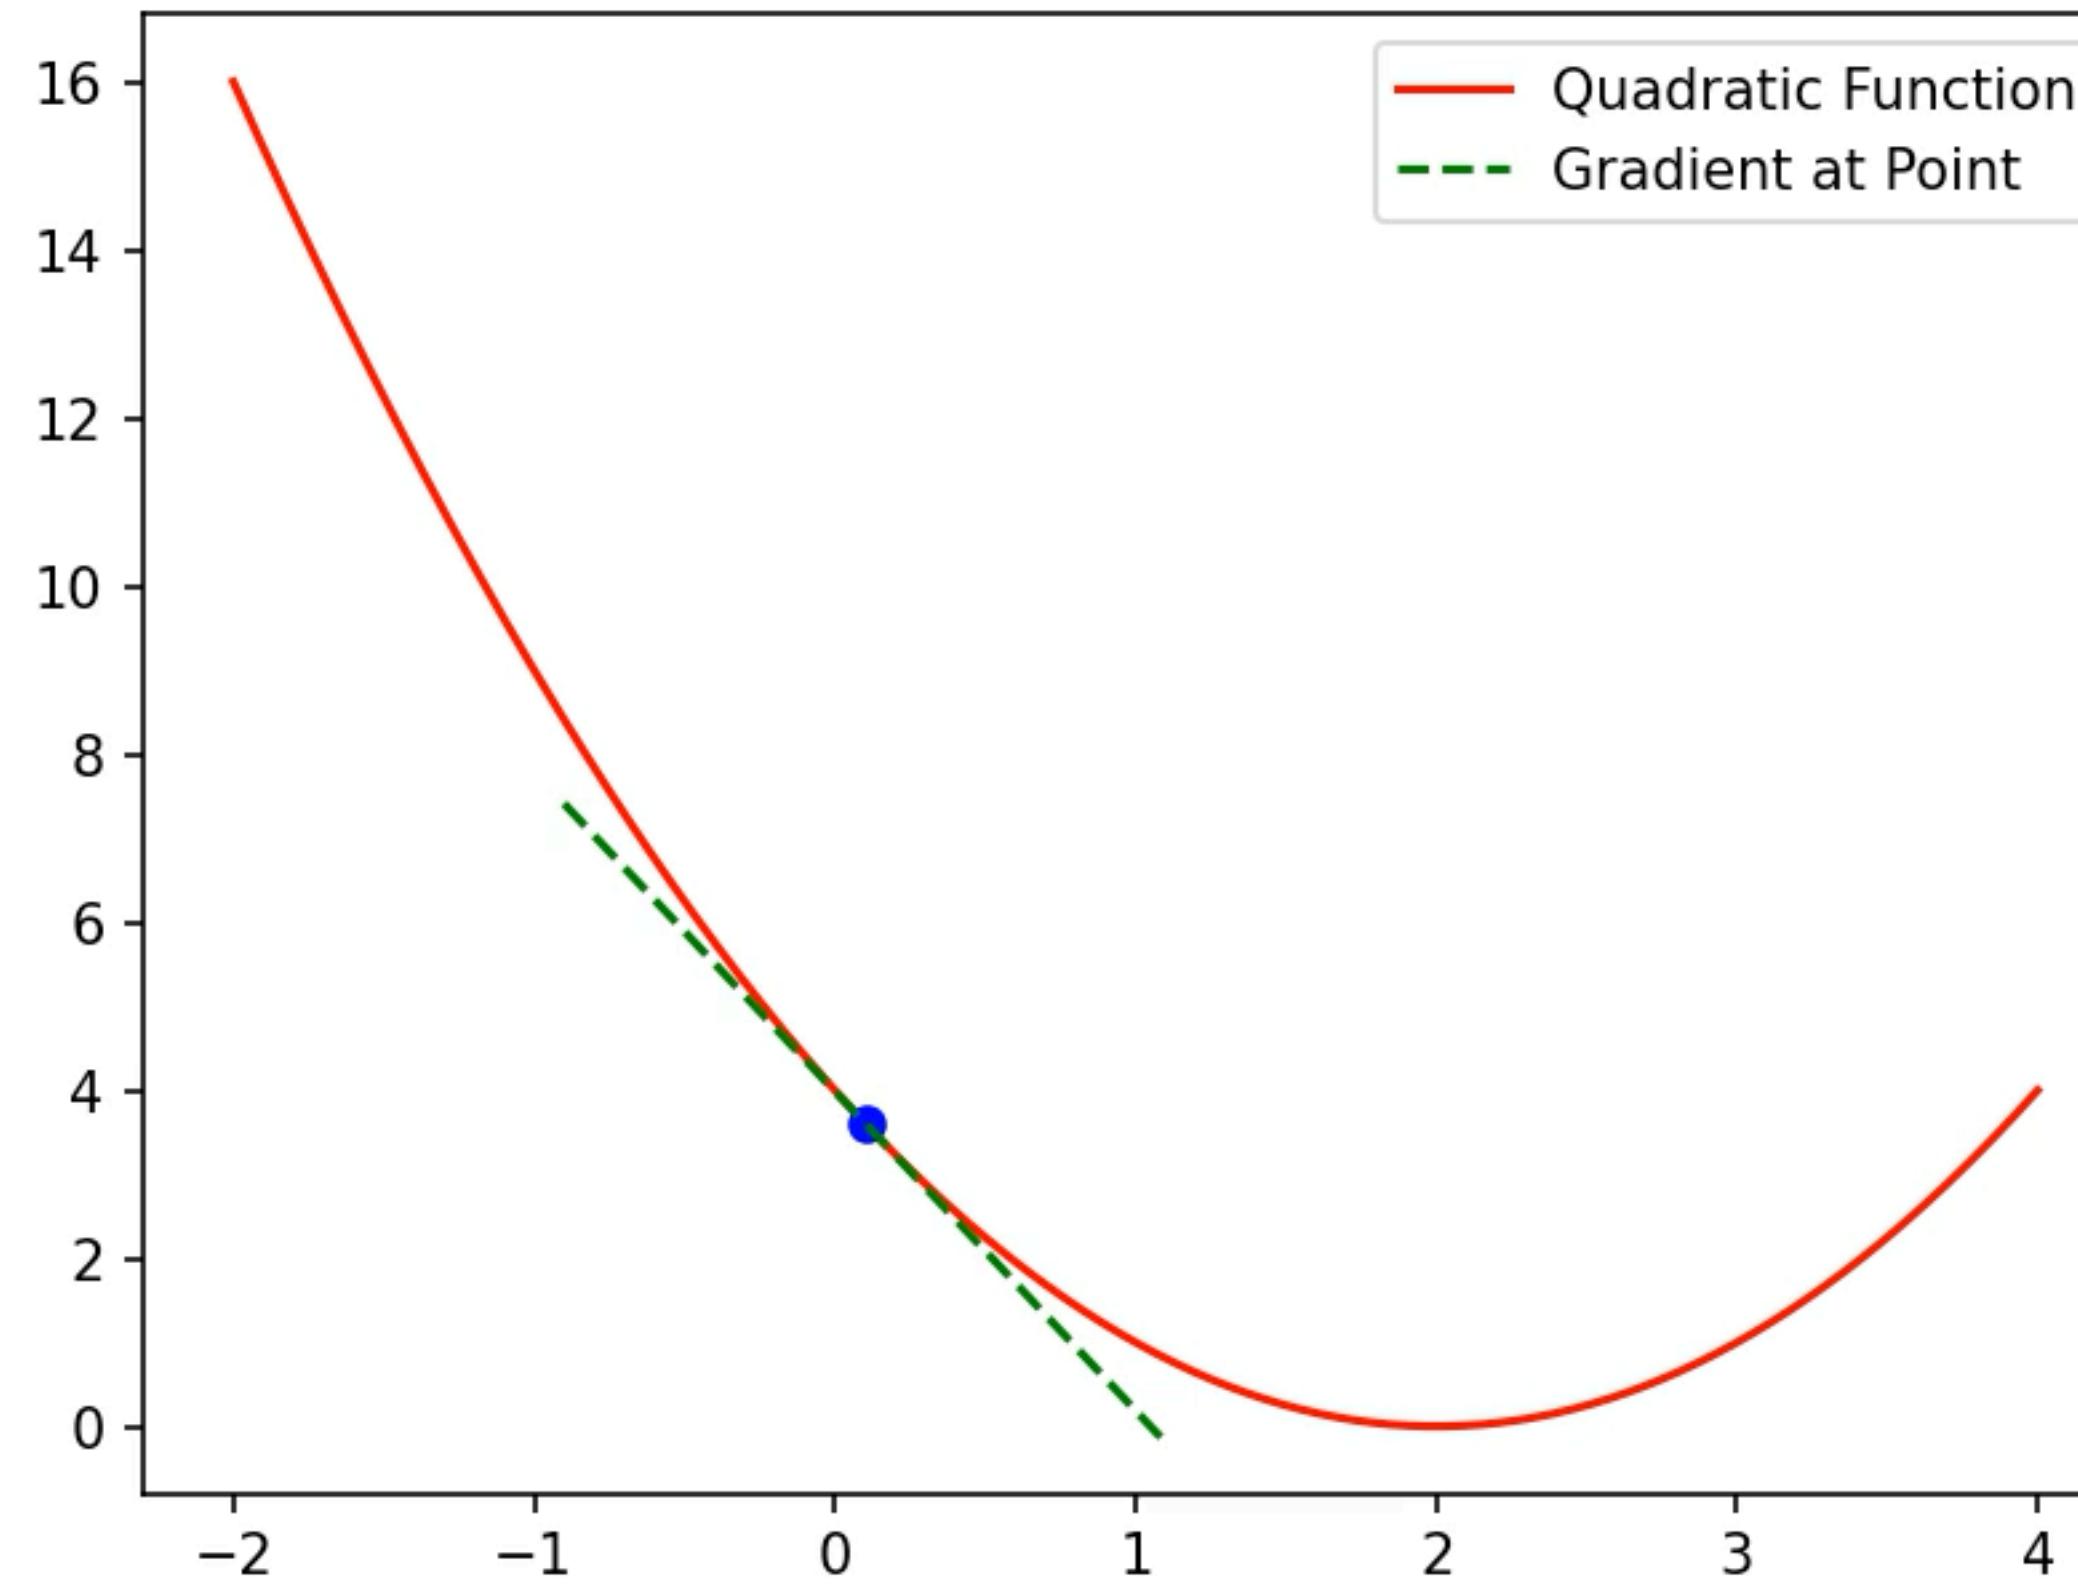
$$\mathbf{x}_{i+1} = \mathbf{x}_i - \eta \nabla f(\mathbf{x}_i)$$

$$\Delta \mathbf{x} = -\eta \nabla f(\mathbf{x})$$

$$\Delta \mathbf{x} = \mathbf{x}_{i+1} - \mathbf{x}_i$$

# OPTIMISATION

## GRADIENT DESCENT - FINDS A LOCAL MINIMUM



$$f'(x) = \lim_{\Delta x \rightarrow 0} \frac{f(x + \Delta x) - f(x)}{\Delta x}$$

$$\Delta \mathbf{x} = -\eta \nabla f(\mathbf{x})$$

$$\Delta \mathbf{x} = \mathbf{x}_{i+1} - \mathbf{x}_i$$

# OPTIMISATION

## IMPROVING GRADIENT DESCENT: KEY STRATEGIES

- Different starting point (initialisation of parameters) may lead to a different solution.
- Exploit noise to get out of shallow local minima.
- Momentum may also help.
- Scaling small gradients.
- Adaptive learning rate.

# RL THROUGH THE LENS OF OPTIMISATION



DALL·E

# RL THROUGH THE LENS OF OPTIMISATION

## REWARD MAXIMISATION

$$G_t = R_{t+1} + R_{t+2} + R_{t+3} + \dots + R_{t+N}$$

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 R_{t+4} + \dots$$

$$0 \leq \gamma < 1$$

# BREAKING DOWN RL TO ITS INGREDIENTS

## LEARN VALUES, CHOOSE ACTIONS

$s_t$

$a_t$

$G_t$

I am in state  $s$  and chose action  $a$ .

I would like to learn the total future rewards  $G$  resulting from that action.

Knowing the expected future reward will allow me to choose a “good” action.

# BREAKING DOWN RL TO ITS INGREDIENTS

## LEARN VALUES, CHOOSE ACTIONS

How do we learn the value of the state-actions?

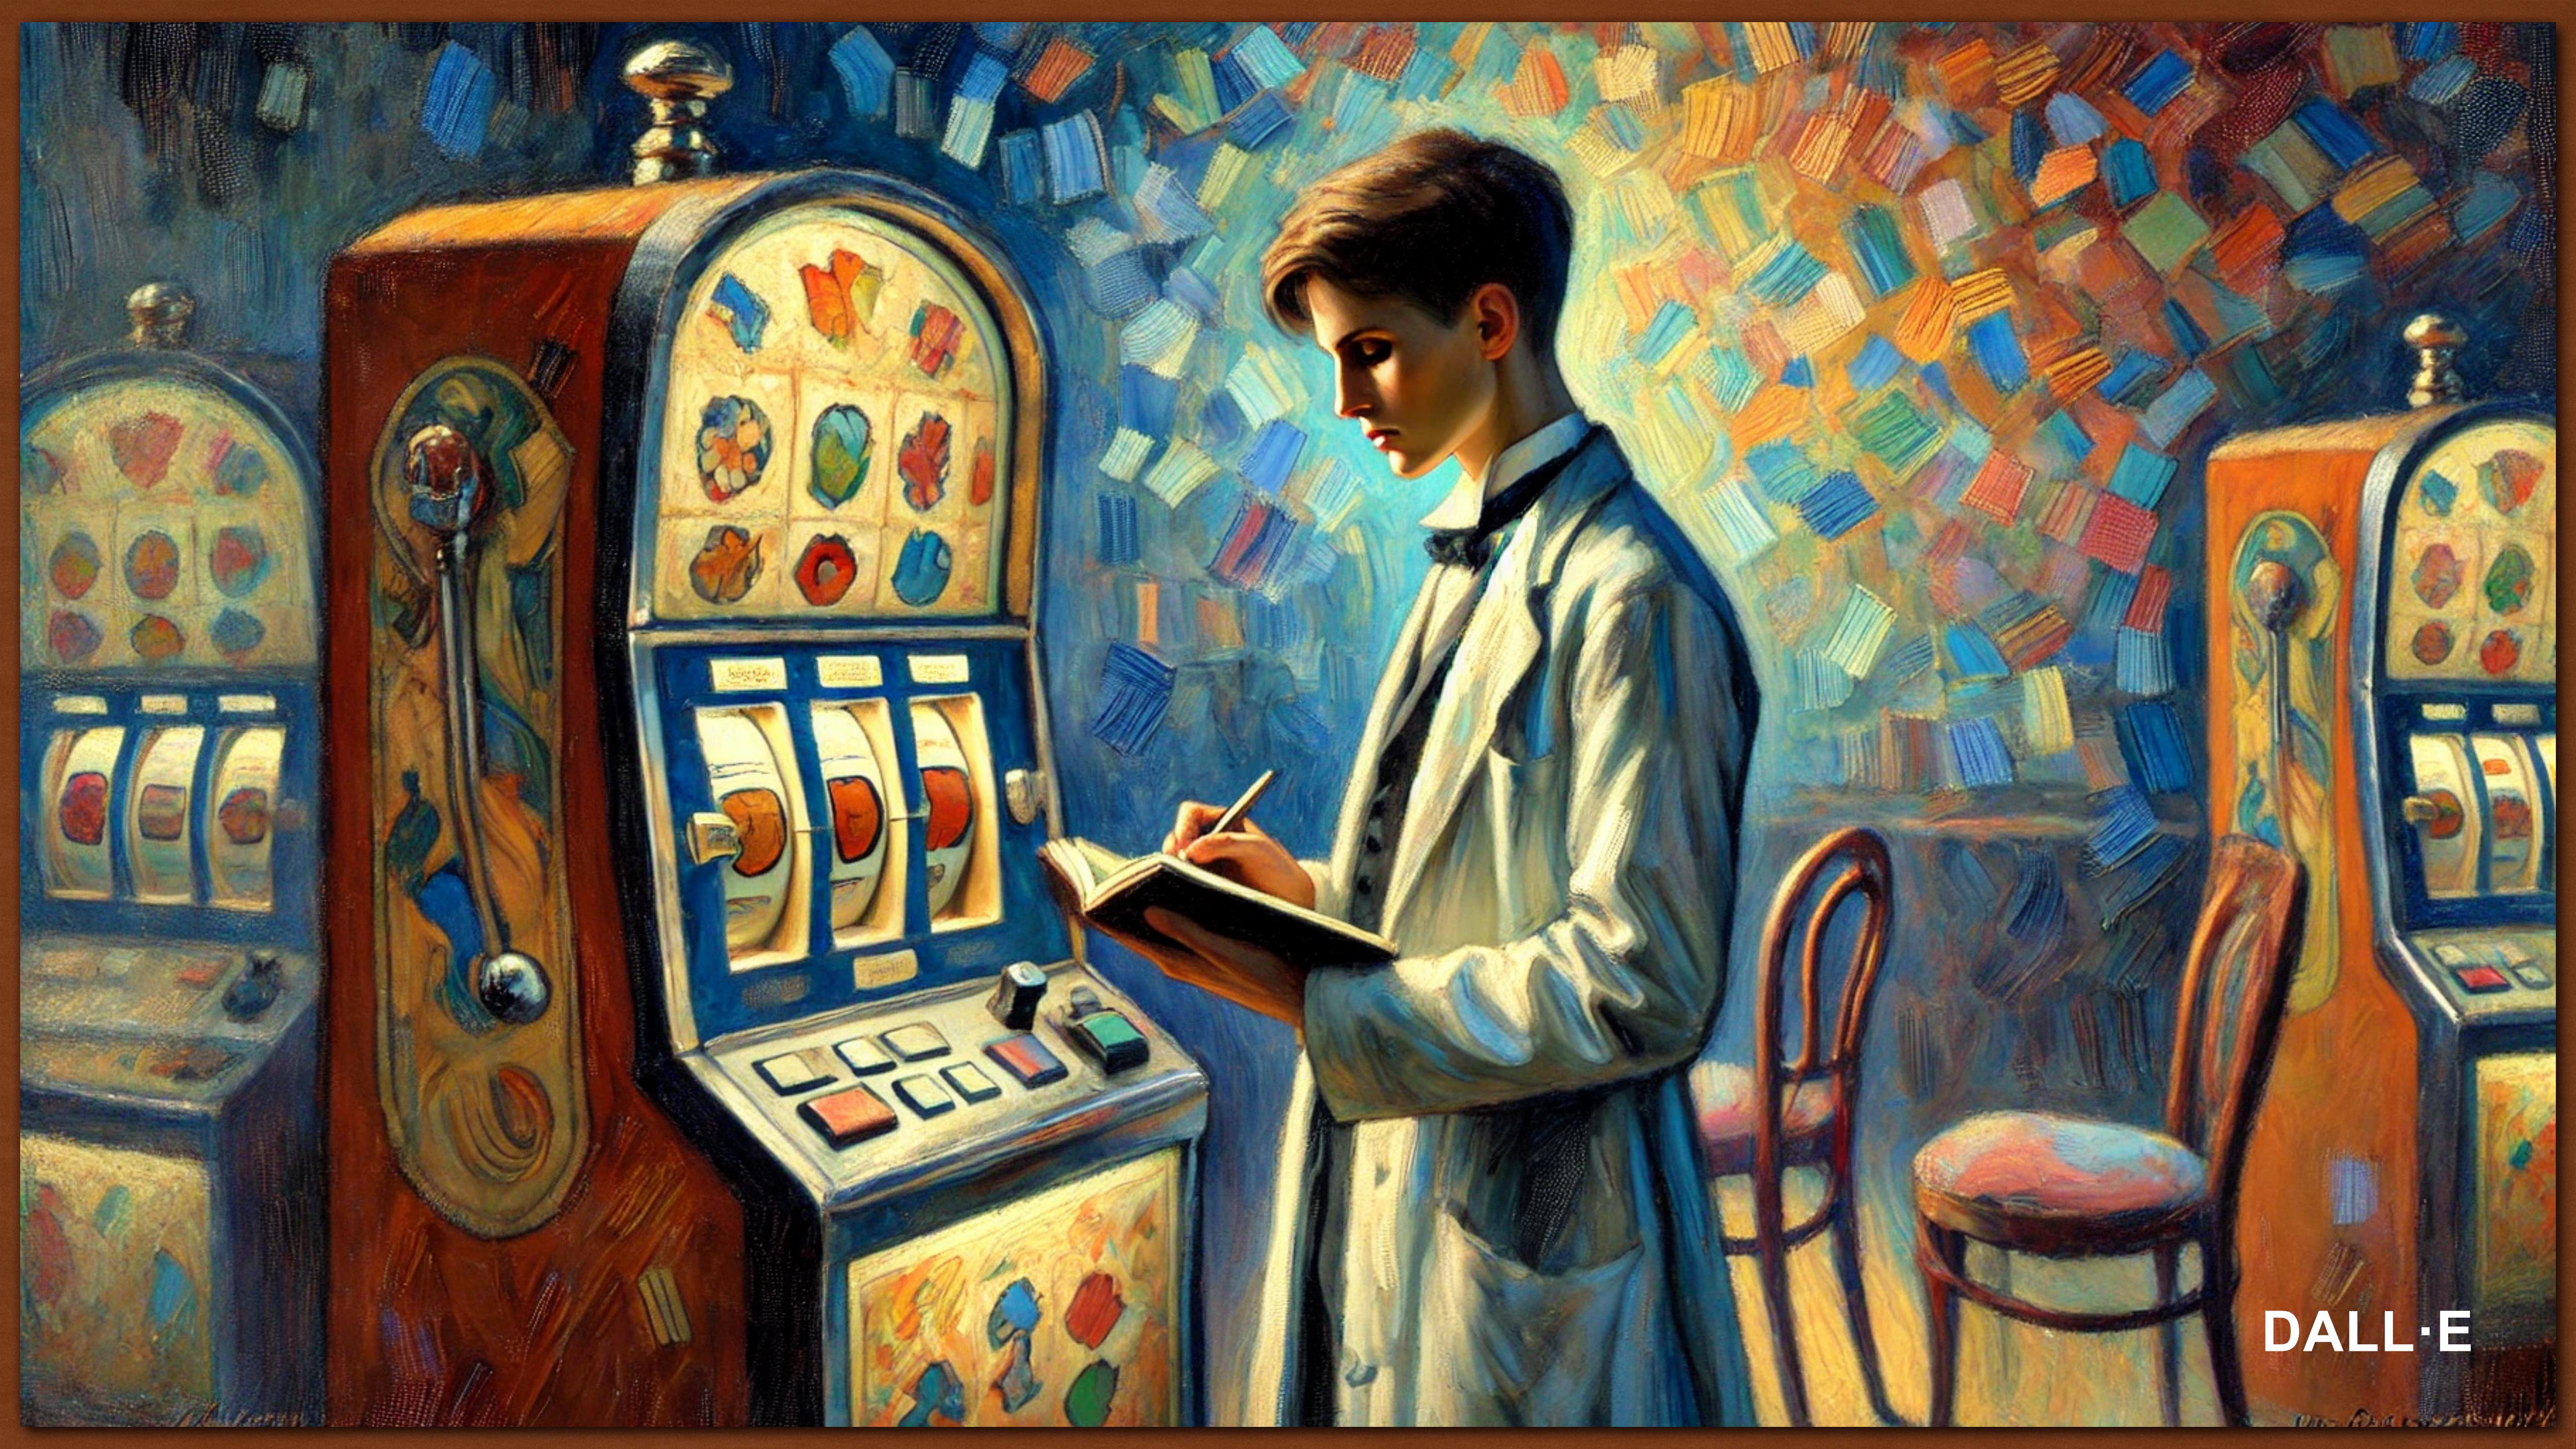
How do we pick an action (aka policy)?

# **RL THROUGH THE LENS OF OPTIMISATION**

## **MAXIMISE R - MINIMISE PREDICTION ERROR**

**Maximise Reward.**

**Minimise prediction error for the expected future reward.**



DALL·E

# IMMEDIATE REWARDS BANDITS

$$G = R(A)$$

No need to denote a state.

Two or more actions.

A “trial” constitutes pulling a lever.

Reward is a consequence of the action of trial n.

# IMMEDIATE REWARDS BANDITS

$q^*$ : true expected return

$$q^*(a) \doteq E[G \mid A_n = a] = E[R \mid A_n = a]$$

The expectation is calculated across trials.

$$E[R \mid A_n = a] = \sum_i r_i P(R = r_i \mid A_n = a)$$

Alternatively, use an empirical estimate:

$$E[R \mid A_n = a] \approx \frac{1}{N(a)} \sum_{n=1}^N R_n \mathbf{1}_{\{A_n=a\}}$$



DALL·E

# IMMEDIATE REWARDS

## BANDITS

$$q^*(a) \doteq E[R | A_n = a]$$

**Q-value : estimate of the expected return**

$$Q(a) \approx q^*(a)$$

**Idea: construct a loss function of Q  
with minimum at:**

$$Q(a) = E[R | A_n = a]$$

**Hint: think quadratic**



DALL·E

# IMMEDIATE REWARDS

## BATCH LOSS FUNCTION

**Minimise:**  $E[(Q(a) - R)^2 \mid A_n = a]$

$$L(a) = \frac{1}{2N(a)} \sum_{n=1}^N (Q(a) - R_n)^2 \mathbf{1}_{A_n=a} \quad N = \sum_a N(a)$$

$\mathbf{1}_{A_n=a} = \begin{cases} 1 & \text{if } A_n = a \\ 0 & \text{otherwise.} \end{cases}$  **Indicator function**

$$N(a) = \sum_{n=1}^N \mathbf{1}_{A_n=a} \quad \mathbf{R}_n \text{ is the reward at trial n}$$



DALL·E

# IMMEDIATE REWARDS

## UPDATE RULE VIA DIRECT MINIMISATION

$$\frac{dL(a)}{dQ(a)} = \frac{d}{dQ(a)} \left( \frac{1}{2N(a)} \sum_{n=1}^N (Q(a) - R_n)^2 \mathbf{1}_{A_n=a} \right)$$

**Because of the convexity of the loss function with respect to Q, we can set directly the derivative to 0 to find the minimum.**

$$\frac{dL(a)}{dQ(a)} = \frac{1}{N(a)} \sum_{n=1}^N (Q(a) - R_n) \mathbf{1}_{A_n=a} = 0$$

**This leads to:**  $Q(a) = \frac{1}{N(a)} \sum_{n=1}^N R_n \mathbf{1}_{A_n=a} \approx E[R \mid A_n = a] = q^*$  **Stationary**

# IMMEDIATE REWARDS

## MULTI-ARM LOSS FUNCTION

$$L = \sum_{a=1}^{|A|} L(a) = \sum_{a=1}^{|A|} \left[ \frac{1}{2N(a)} \sum_{n=1}^N (Q(a) - R_n)^2 \mathbf{1}_{A_n=a} \right]$$

In this case, we would need to take partial derivatives

$$\frac{\partial L}{\partial Q(a)}$$

and end up with the same result.

# IMMEDIATE REWARDS

## GRADIENT BATCH UPDATE

$$\frac{dL(a)}{dQ(a)} = \frac{1}{N(a)} \sum_{n=1}^N (Q(a) - R_n) \mathbf{1}_{A_n=a}$$

We can also write a gradient rule:

$$\Delta Q(a) = -\eta N(a) \frac{dL(a)}{dQ(a)} = -\eta \sum_{n=1}^N (Q(a) - R_n) \mathbf{1}_{A_n=a}$$

**η: learning rate**

Batch rule, first collect observations (trial data), then update.

# IMMEDIATE REWARDS

## GRADIENT BATCH UPDATE MATCHES DIRECT OPTIMISATION

$$\Delta Q(a) = -\eta \sum_{n=1}^N (Q(a) - R_n) \mathbf{1}_{A_n=a}$$

At convergence  $\Delta Q(a)=0$ , hence:

$$Q(a) = \frac{1}{N(a)} \sum_{n=1}^N R_n \mathbf{1}_{A_n=a} \quad \text{exactly as calculated earlier}$$

# IMMEDIATE REWARDS

## ONLINE GRADIENT RULE

$$\Delta Q(a) = -\eta \frac{dL(a)}{dQ(a)} = -\eta \sum_{n=1}^N (Q(a) - R) \mathbf{1}_{A_t=a}$$

We now update our estimate after each trial (sample) known as online learning. We then write:

$$\Delta Q(a) = -\eta (Q(a) - R_n) \mathbf{1}_{A_n=a}$$

i.e. for trial 1 to N, if action  $a$  is taken, we update its Q value by:

$$Q(a) = Q(a) - \eta (Q(a) - R_n)$$

# IMMEDIATE REWARDS

## WHEN TO USE THE ONLINE RULE

Lets assume a stationary environment, introduce an index k on the trials where action a is chosen and set  $\eta=1/(k+1)$ .

$$Q_{k+1}(a) = Q_k(a) - \frac{1}{k+1}(Q_k(a) - R_{k+1}).$$

We can use induction to show that it converges to the numerical average, just like the batch version.

Also, for fixed learning rate  $\eta$ , we can easily show that it takes the form:

$$Q_{k+1}(a) = (1 - \eta)Q_k(a) + \eta R_{k+1} \quad \text{Exponential average (forgetful), non-stationary}$$

# POLICIES

## THE EXPLORATION EXPLOITATION DILEMMA



DALL·E

# POLICIES

## THE EXPLORATION EXPLOITATION DILEMMA

**Greedy**  $a_n = \operatorname{argmax}_a Q(a)$

**Optimistic Greedy:** initialise Q-values unrealistically high

**Epsilon-Greedy :** explore with probability epsilon, greedy otherwise

**Softmax:**  $P(a) = \frac{e^{Q(a)/\tau}}{\sum_b e^{Q(b)/\tau}}$

What happens if  $\tau$  grows very large or tends to 0?

# IMMEDIATE REWARDS

## BANDITS

- Initialise Q-values, and count  $N(a)$  for all actions  $a$
- For each trial  $n$ 
  - Select an action  $a$  using policy
  - Increase the count  $N(a)$
  - Take action  $a$  and observe reward  $R$
  - Update  $Q(a)$ 
    - For a stationary environment:  $Q(a) = Q(a) + (R - Q(a))/N(a)$
    - For non-stationary environment:  $Q(a) = Q(a) + \eta(Q(a) - R)$



# IMMEDIATE REWARDS BANDITS



DALL·E

# IMMEDIATE REWARDS

## BANDITS



# BANDITS

## SUMMARY

- Optimisation
- Immediate Rewards (bandits)
- Batch update rule
- Online update rule
- Bandit algorithm for stationary and non-stationary environments

**THANK YOU!**