

COM 3240

REINFORCEMENT LEARNING



OPTIMISATION

GRADIENT DESCENT

Desirable: A system (an agent) that performs well a task

The system has parameters that we need to select appropriately (optimise)

$f(x_1, x_2, \dots, x_n)$
parameters

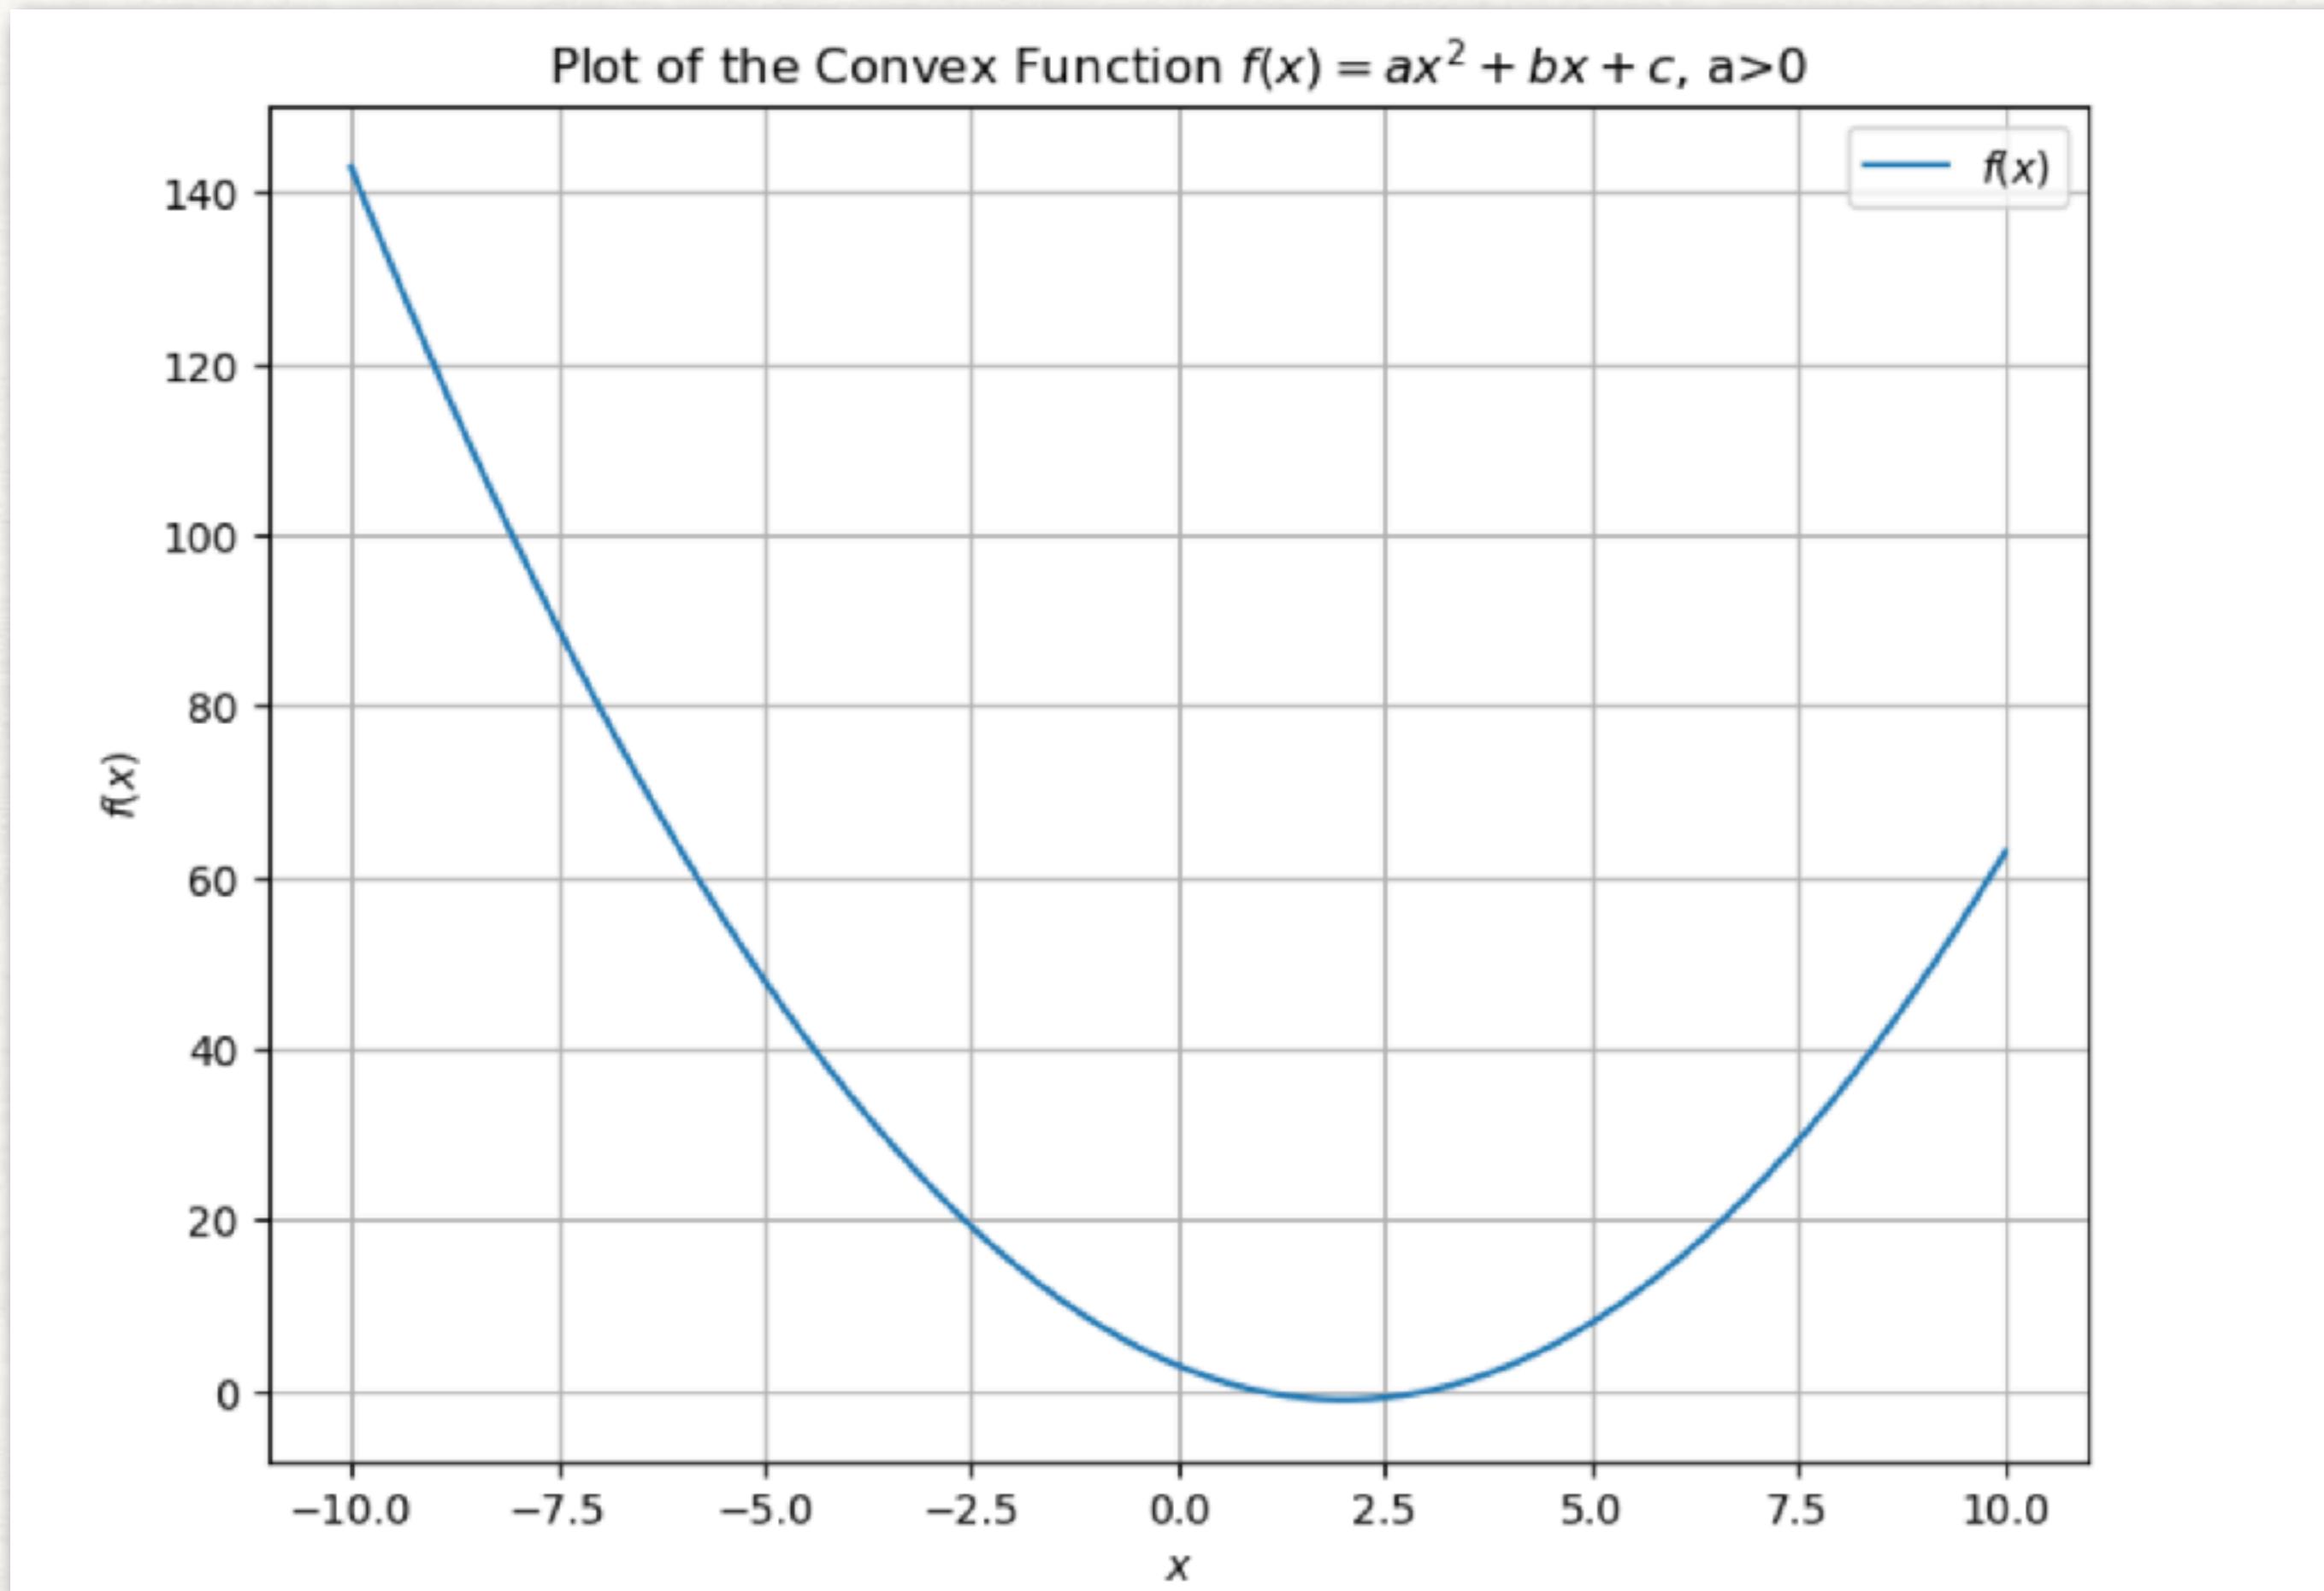
f System's performance: maximise
System's error : minimise



DALL·E

OPTIMISATION

SINGLE (GLOBAL) MINIMUM




$$\frac{d}{dx}(ax^2 + bx + c) = 2ax + b$$

$$2ax + b = 0$$

OPTIMISATION

SINGLE (GLOBAL) MINIMUM

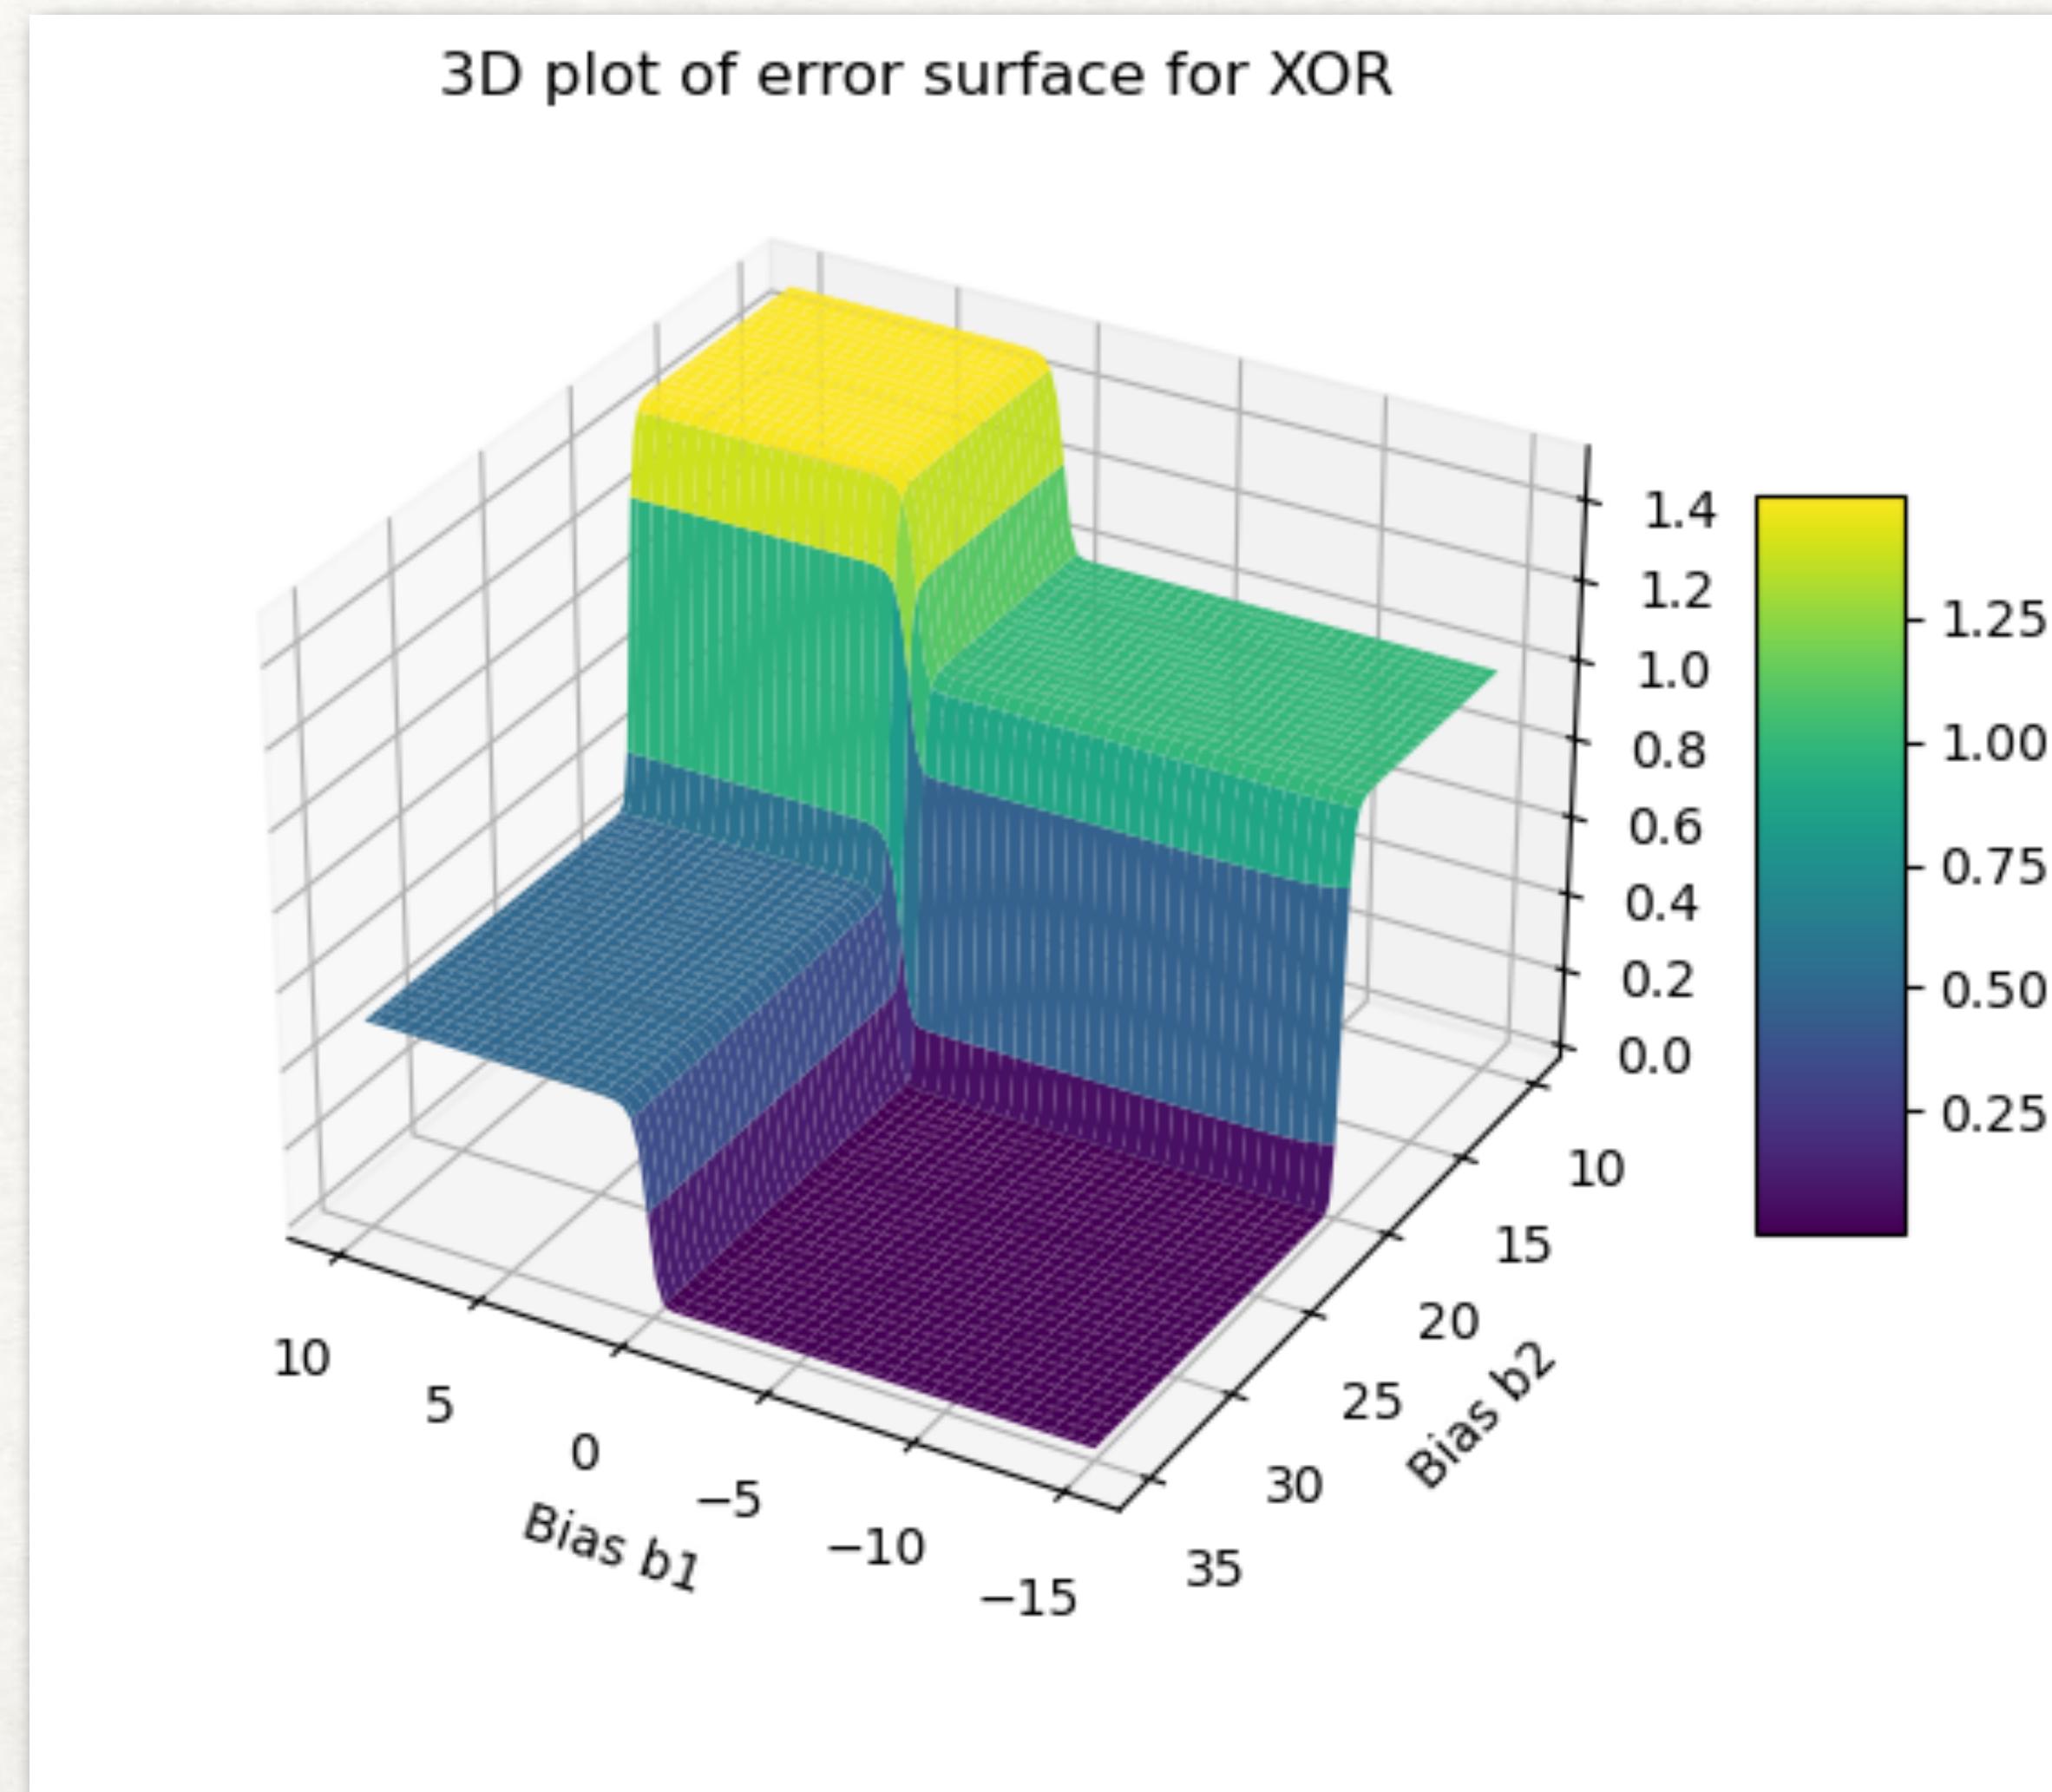


$$\nabla f(\mathbf{x}) = \left[\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2} \right] = 0$$

But this is not the case for all machine learning techniques, including multi-layer and deep neural networks

OPTIMISATION - GRADIENT DESCENT

MULTIPLE LOCAL MINIMA



Simple ANN example

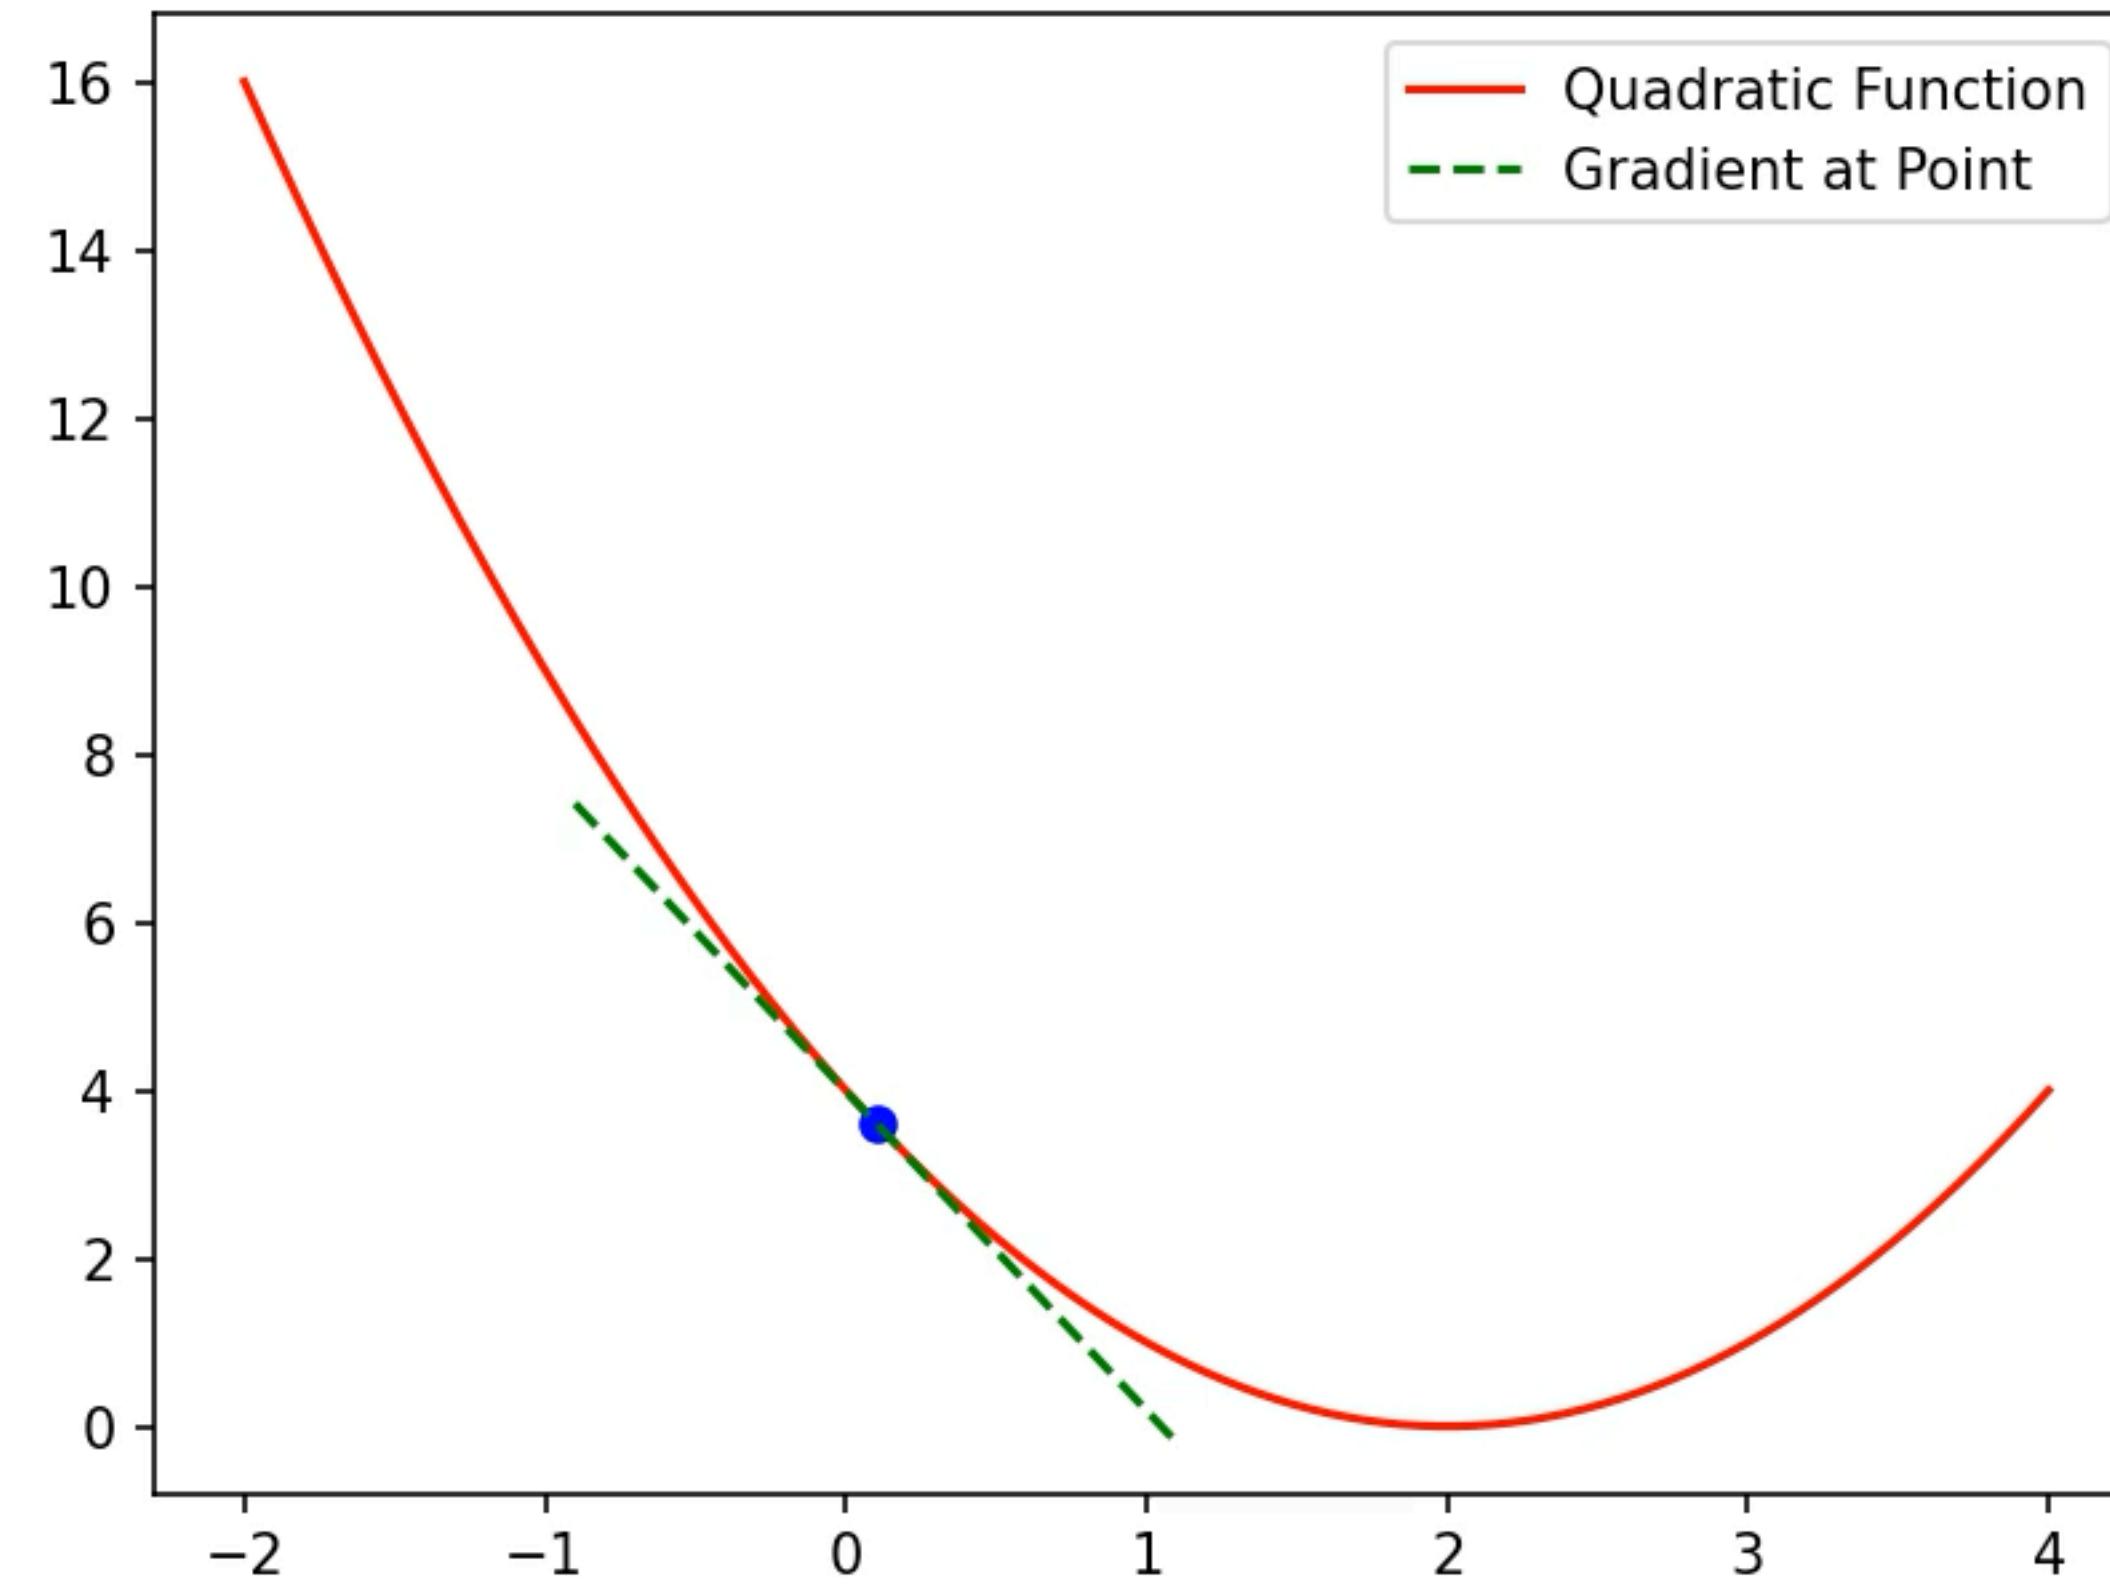
$$\mathbf{x}_{i+1} = \mathbf{x}_i - \eta \nabla f(\mathbf{x}_i)$$

$$\Delta \mathbf{x} = -\eta \nabla f(\mathbf{x})$$

$$\Delta \mathbf{x} = \mathbf{x}_{i+1} - \mathbf{x}_i$$

OPTIMISATION

GRADIENT DESCENT - FINDS A LOCAL MINIMUM



$$f'(x) = \lim_{\Delta x \rightarrow 0} \frac{f(x + \Delta x) - f(x)}{\Delta x}$$

$$\Delta \mathbf{x} = -\eta \nabla f(\mathbf{x})$$

$$\Delta \mathbf{x} = \mathbf{x}_{i+1} - \mathbf{x}_i$$

OPTIMISATION

IMPROVING GRADIENT DESCENT: KEY STRATEGIES

- Different starting point (initialisation of parameters) may lead to a different solution.
- Exploit noise to get out of shallow local minima.
- Momentum may also help.
- Scaling small gradients.
- Adaptive learning rate.

RL THROUGH THE LENS OF OPTIMISATION



RL THROUGH THE LENS OF OPTIMISATION

REWARD MAXIMISATION

$$G_t = R_{t+1} + R_{t+2} + R_{t+3} + \dots + R_{t+N}$$

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 R_{t+4} + \dots$$

$$0 \leq \gamma < 1$$

BREAKING DOWN RL TO ITS INGREDIENTS

LEARN VALUES, CHOOSE ACTIONS



G_t

I am in state s and chose action a .

I would like to learn the total future rewards G resulting from that action.

Knowing the expected future reward will allow me to choose a “good” action.

BREAKING DOWN RL TO ITS INGREDIENTS

LEARN VALUES, CHOOSE ACTIONS

How do we learn the value of the state-actions?

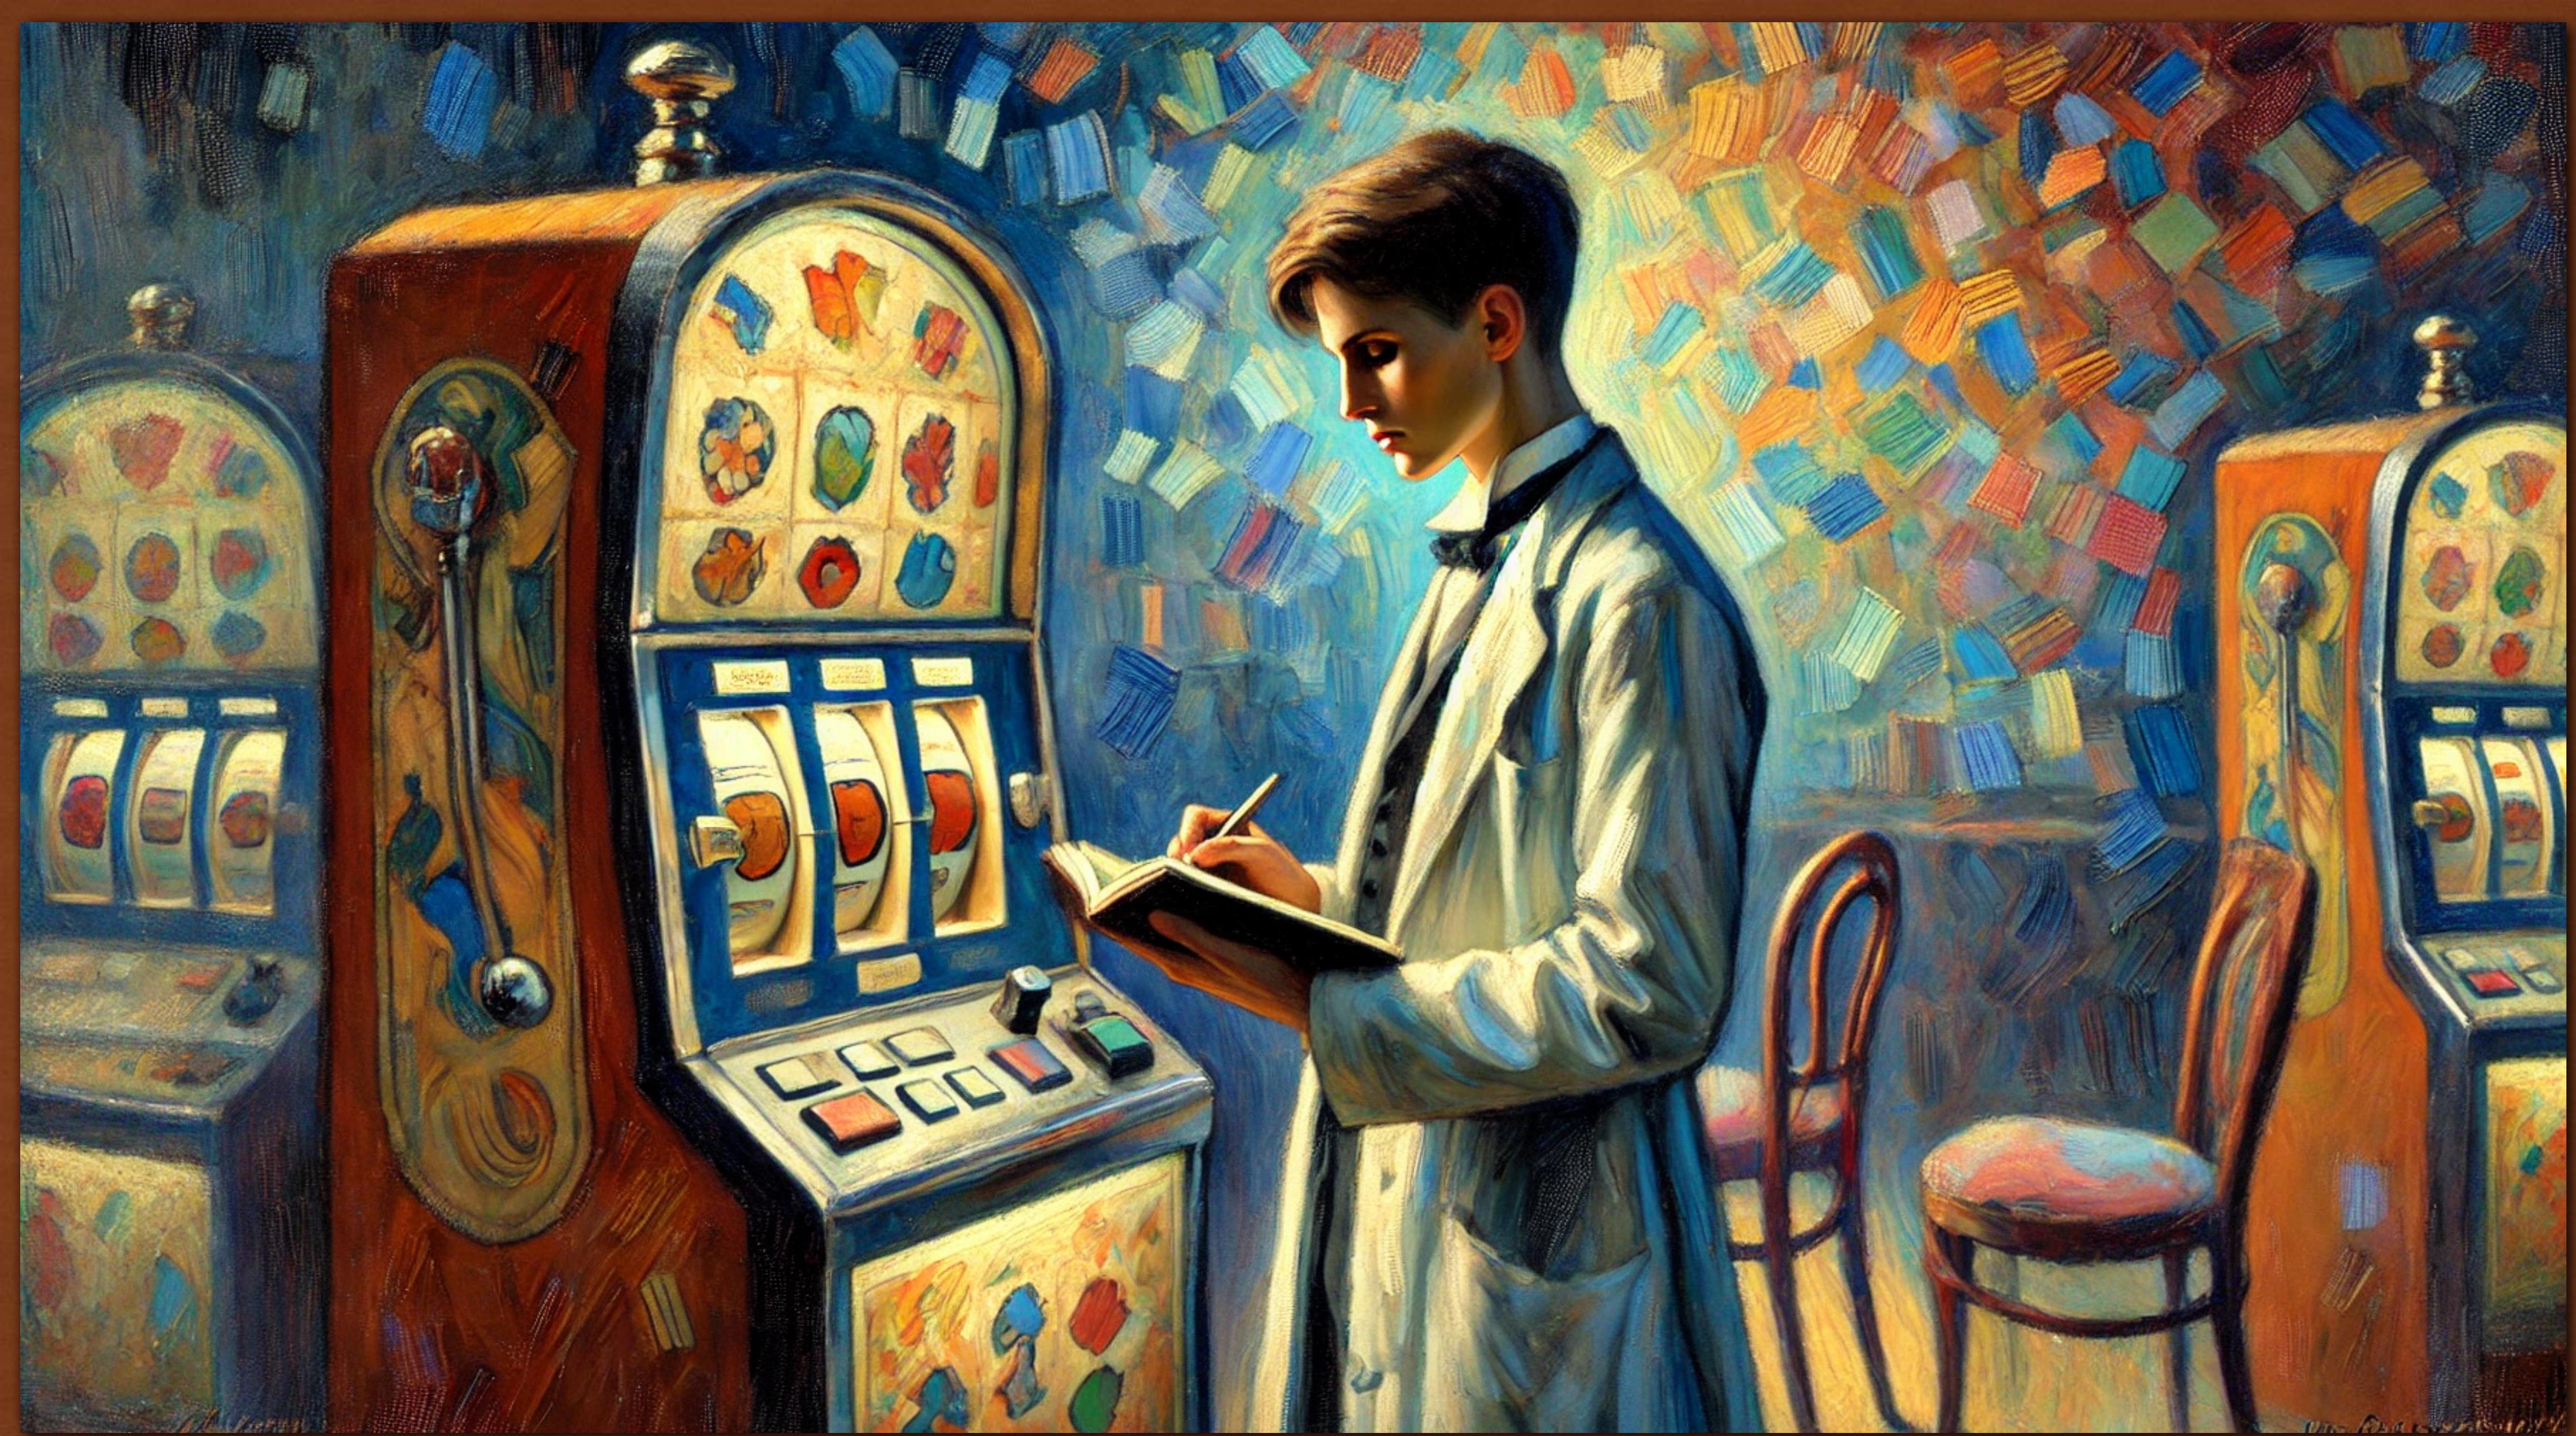
How do we pick an action (aka policy)?

RL THROUGH THE LENS OF OPTIMISATION

MAXIMISE R - MINIMISE PREDICTION ERROR

Maximise Reward.

Minimise prediction error for the expected future reward.



IMMEDIATE REWARDS

BANDITS

$$G_t = R(A_t)$$

No need to denote a state.

Two or more actions.

Time t here is an index on the “trial”, which constitutes pulling a lever.

Reward is a consequence of the action of trial t.

IMMEDIATE REWARDS BANDITS

q^* : estimate of true expected reward

$$\begin{aligned} q^*(a) &\doteq E[G_t | A_t = a] \\ &\doteq E[R(A_t) | A_t = a] = E[R(A_t = a)] \end{aligned}$$

The expectation is calculated across trials.

$$E[R(A_t = a)] = \sum_i r_i P(R(A_t = a) = r_i)$$

Rewards and reward distributions
are unknown: $r_i, P(R = r_i)$



DALL·E

IMMEDIATE REWARDS BANDITS

$$q^*(a) \doteq E[G_t | A_t = a] = E[R(A_t = a)]$$

Q-value : estimate of the expected return

$$Q(a) \approx q^*(a)$$

In RL we obtain an estimate via sampling

Ideally: $\frac{E[(Q(a) - R(A_t = a))^2]}{2} = 0$



DALL·E

IMMEDIATE REWARDS BANDITS

Minimise: $\frac{E[(Q(a) - R(A_t = a))^2]}{2}$

$$L(a) = \frac{1}{2T_a} \sum_{t=1}^T (Q(a) - R)^2 \mathbf{1}_{A_t=a}$$

$\mathbf{1}_{A_t=a} = \begin{cases} 1 & \text{if } A_t = a \\ 0 & \text{otherwise.} \end{cases}$ Indicator function

$$T(a) = \sum_{t=1}^T \mathbf{1}_{A_t=a} \quad R=R(t) \text{ is the reward at trial t}$$



DALL·E

IMMEDIATE REWARDS BANDITS

$$\frac{dL(a)}{dQ} = \frac{1}{T(a)} \sum_{t=1}^T (Q(a) - R) \mathbf{1}_{A_t=a}$$

Because of the convexity of the loss function with respect to Q, we can set directly the derivative to 0 to find the minimum.

$$\frac{1}{T(a)} \sum_{t=1}^T (Q(a) - R) \mathbf{1}_{A_t=a} = 0$$

This leads to: $Q(a) = \frac{1}{T(a)} \sum_{t=1}^T R \mathbf{1}_{A_t=a} \approx E[R(A_t = a)] = q^*$ **Stationary**

IMMEDIATE REWARDS BANDITS

$$L = \sum_a L(a) = \sum_a \left[\frac{1}{2T(a)} \sum_{t=1}^T (Q(a) - R)^2 \mathbf{1}_{A_t=a} \right] \quad T = \sum_a T(a)$$

In this case, we would need to take partial derivatives

$$\frac{\partial L}{\partial Q(a)}$$

and end up with the same result.

IMMEDIATE REWARDS

BATCH AND ONLINE UPDATES

$$\frac{dL(a)}{dQ} = \frac{1}{T(a)} \sum_{t=1}^T (Q(a) - R) \mathbf{1}_{A_t=a}$$

We can also write a gradient rule:

$$\Delta Q(a) = -\eta T(a) \frac{dL(a)}{dQ} = -\eta \sum_{t=1}^T (Q(a) - R) \mathbf{1}_{A_t=a}$$

η : learning rate

Batch rule, first collect observations (trial data), then update.

IMMEDIATE REWARDS

BATCH AND ONLINE UPDATES

$$\Delta Q(a) = -\eta \frac{dL(a)}{dQ} = -\eta \sum_{t=1}^T (Q(a) - R) \mathbf{1}_{A_t=a}$$

At convergence $\Delta Q(a)=0$, hence:

$$Q(a) = \frac{1}{T(a)} \sum_{t=1}^T R \mathbf{1}_{A_t=a}$$

exactly as calculated earlier

IMMEDIATE REWARDS

BATCH AND ONLINE UPDATES

$$\Delta Q(a) = -\eta \frac{dL(a)}{dQ} = -\eta \sum_{t=1}^T (Q(a) - R) \mathbf{1}_{A_t=a}$$

Note a: action,
α: learning rate!

We can suppress the sum. This suggests we now update our estimate after each trial (sample) known as online learning. We then write:

$$\Delta Q(a) = -\eta (Q(a) - R) \mathbf{1}_{A_t=a}$$

i.e. for trial 1 to T, if action a is taken, we update its Q value by:

$$Q(a) = Q(a) - \eta (Q(a) - R)$$

IMMEDIATE REWARDS

THE ONLINE RULE MAKES SENSE

Lets assume a stationary environment, introduce an index i on the trials where action a is chosen and set $\eta=1/(i+1)$.

$$Q_{i+1}(a) = Q_i(a) + \frac{1}{i+1} (R_{i+1} - Q_i(a)).$$

We can use induction to show that it converges to the numerical average.

Also, for fixed learning rate η , we can easily show that it takes the form:

$$Q_{i+1}(a) = (1 - \eta)Q_i(a) + \eta R_{i+1}$$

Exponential average (forgetful), non-stationary

POLICIES

THE EXPLORATION EXPLOITATION DILEMMA



DALL·E

POLICIES

THE EXPLORATION EXPLOITATION DILEMMA

Greedy $a_t = \operatorname{argmax}_a Q_t(a)$

Optimistic Greedy: initialise Q-values unrealistically high

Epsilon-Greedy : explore with probability epsilon, greedy otherwise

Softmax: $P(a) = \frac{e^{Q_t(a)/\tau}}{\sum_b e^{Q_t(b)/\tau}}$

What happens if τ grows very large or tends to 0?

IMMEDIATE REWARDS

BANDITS

- Initialise Q-values, and count $T(a)$ for all actions a
- For each time step
 - Select an action a using policy
 - Increase the count $T(a)$
 - Take action a and observe reward R
 - Update $Q(a)$
 - For a stationary environment: $Q(a) = Q(a) + (R - Q(a))/T(a)$
 - For non-stationary environment: $Q(a) = Q(a) + \eta(R - Q(a))$


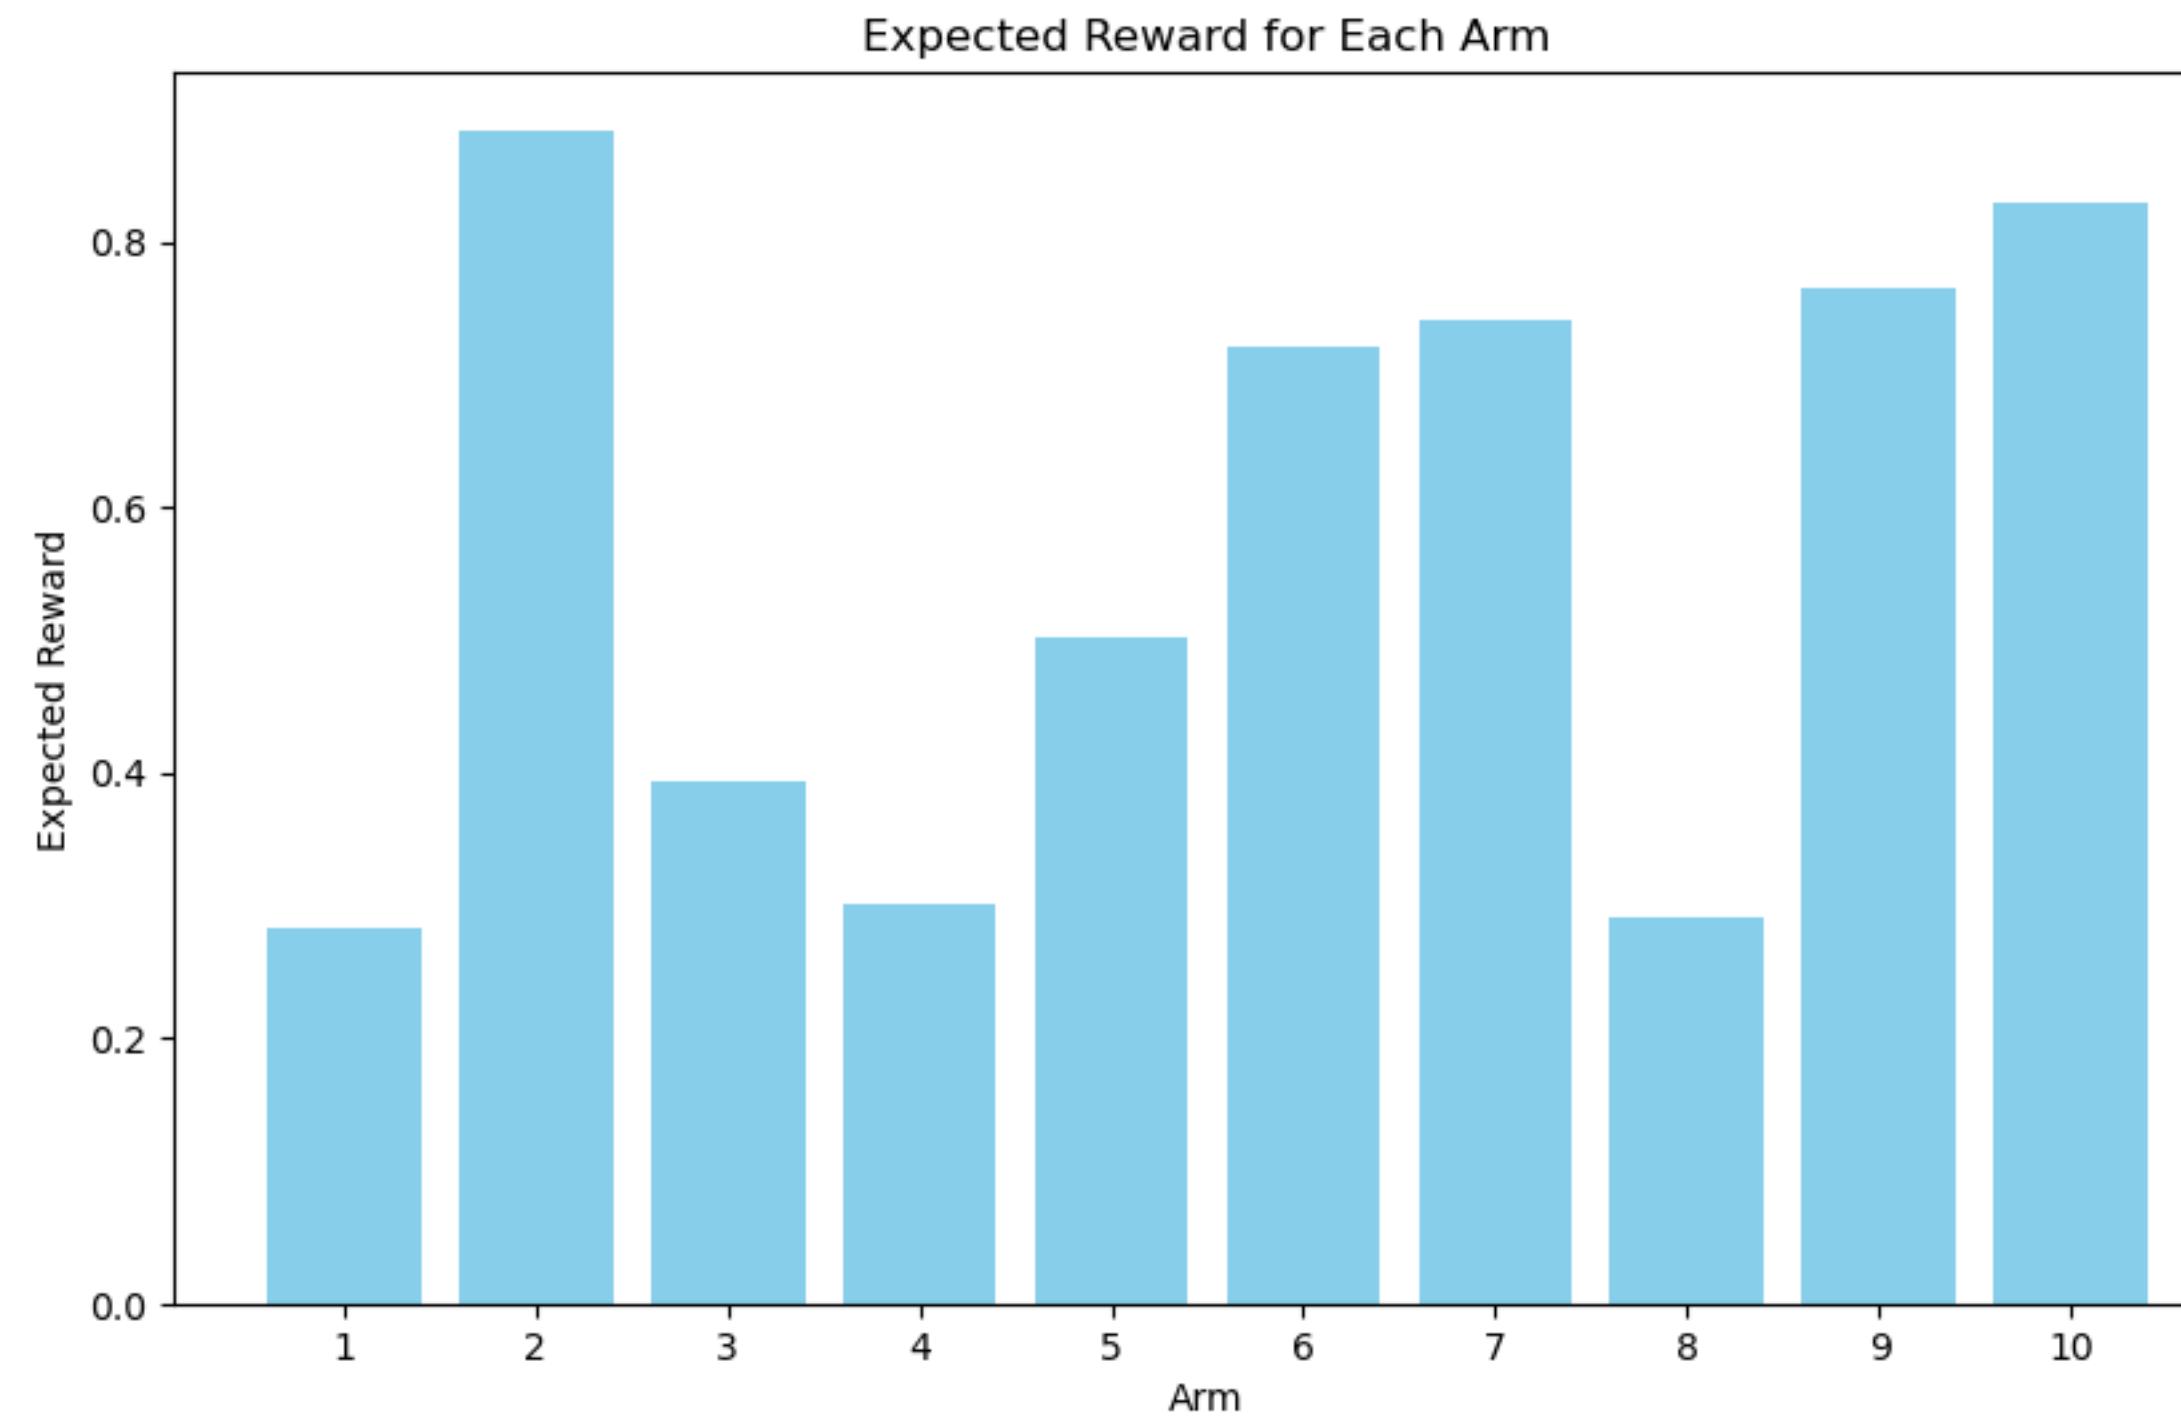
IMMEDIATE REWARDS BANDITS



DALL·E

IMMEDIATE REWARDS

BANDITS



BANDITS

SUMMARY

- Optimisation
- Immediate Rewards
- Batch Learning
- Online Learning
- Bandit algorithm for stationary and non-stationary environments

THANK YOU!